

Databasesystems 2

Forum: <https://forum-db.informatik.uni-tuebingen.de/c/ss18-db2>

Assignment 8 (19.06.2018)


Submission: Tuesday, 26.06.2018, 10:00 AM

Please note that students currently have the opportunity to **evaluate lectures**. Please help us to improve **your** courses by providing precious feedback. Check your Mailbox now and participate **today**.

1. [8 Points] **B⁺Tree - Variation**

Consider the B⁺Tree of order $o = 2$ in Figure 1 where the Insert- and Delete-operations do **not** implement *redistribution*.

- Find five *additional* key values a, \dots, e such that the following holds: If you insert the entries in the order a, \dots, e and then delete them in the reverse order e, \dots, a , the resulting tree is the **same** as in Figure 1.
- Find five *additional* key values a, \dots, e such that the following holds: If you insert the entries in the order a, \dots, e and then delete them in the reverse order e, \dots, a , the resulting tree **differs** from the one in Figure 1.

Figure 1: A B⁺Tree2. [11 Points] **Cluster Factor**

Even if a table is clustered according to an index, subsequent INSERT and UPDATE operations may destroy the table's perfect sort order. Instead tuples may be inserted on any page of the table, possibly reducing the *clustering factor* of the table.

Clustered Tuple: A tuple t of table T is *clustered* w.r.t an index I , if the tuple t' , immediately preceding t in the sequence set of I , is located on the same page of T as t or on a page of T physically located before the page of t . The first tuple in the sequence set of I is always clustered.

Clustering Factor: The *clustering factor* of table T is $\frac{|\text{clustered tuples in } T|}{|\text{tuples in } T|}$

Load file `cluster-factor.sql` to create table `indexed(a int, b text, c numeric(3,2))` as introduced in the lectures. Table `indexed` is equipped with two indexes, `indexed.a` on column `a` and `indexed.c.a` on composite key $\langle c, a \rangle$. The table is clustered according to index `indexed.c.a`.

- Write a query to compute the *clustering factor* of table `indexed` with respect to index `indexed.c.a`.
Hint: Use window function `LAG()`¹ to access the tuple immediately preceding the current tuple with respect to the given sort order.
- Write a query to compute the *clustering factor* of table `indexed` with respect to index `indexed.a`. Explain the difference to the result of 2a, **briefly**.

¹<https://www.postgresql.org/docs/current/static/functions-window.html#FUNCTIONS-WINDOW-TABLE>

3. [11 Points] UB-Trees: B⁺Tree on Z-order values

One possibility to index a table on multiple columns are *composite indexes*. However these do not support both dimensions equally. Instead, the first dimension dominates the order of entries in a composite index.

Consider a table `points(x INT, y INT)` representing points in a two dimensional space:

```
-- A two dimensional space
CREATE TABLE points (x INT, y INT);


-- Populate space with points
INSERT INTO points
  (SELECT x, y
   FROM   generate_series(0,99) AS x, generate_series(0,99) AS y);

-- Create primary key index
CREATE UNIQUE INDEX points_x_y ON points USING btree (x,y);
```

Composite B⁺Tree indexes reflect the spacial locality of points in a two-dimensional space only in the first dimension *x*. Locality regarding *both* dimensions can be achieved with a B⁺Tree index over an expression that combines both dimensions in a so called *Z-order value*:

```
CREATE INDEX points_z_value_x_y ON points USING btree (z_order_value(x,y));
ANALYZE;
```

The Z-order value of a point (x,y) is calculated using the bit-interleaving function provided in `zorder.sql`.² This index reflects locality for both input dimensions:



- Write a SQL query that, given an index `I` on table `points` and an index leaf page number `n`, returns all points that are located on `I`'s leaf page `n`. Use function `bt_page_items(indexname TEXT, pageno INT)` provided by `CREATE EXTENSION pageinspect` to access the *RIDs* of all items on a B⁺Tree page (e.g. `pageno = 7`). Assume that the given page is a leaf page.
- For both indexes, `points_x_y` and `points_z_value_x_y`:
 - Choose an arbitrary leaf page.
(Check that attribute `type` returned by `bt_page_stats(indexname, pageno)` is `'l'`.)
 - Collect the points located on this page using the query of 3a.
 - Illustrate the `x`, `y` location of these points in the two-dimensional space, plotting them on a grid of size 100×100 with *Gnuplot*.
You can access a web-based version of Gnuplot at <http://gnuplot.respawned.com>.³
 - Describe and explain your findings briefly.

²See Wikipedia for more information on Z-order values: https://en.wikipedia.org/wiki/Z-order_curve

³Examples for "Data" and "Plot script" can be found in files `data.txt` and `points.plot`.