

Databasesystems 2

Forum: <https://forum-db.informatik.uni-tuebingen.de/c/ss18-db2>

Assignment 11 (10.07.2018)

Submission: Tuesday, 17.07.2018, 10:00 AM

1. [15 Points] **Implementation of Relational Operators**

The following examples outline the semantics of the *Relational Union* (\cup) and *Left Semi Join* ($\ltimes_{a=b}$) which are operators used in the relational algebra:

- *Union* (\cup): The result contains all rows of R_1 and R_2 without duplicates:

R_1		R_2		R_{res}
A B		A B		A B
2 4		2 4	=	2 4
6 1		3 5		6 1
10 5		4 4		10 5
8 4		10 5		8 4
		6 5		3 5
				4 4
				6 5

- *Left Semi Join* ($\ltimes_{a=b}$): The result contains every row of R_1 that finds at least one join partner in R_2 :

R_1		R_2		R_{res}
A B C		D E		A B C
2 4 2		2 4	=	2 4 2
9 1 3		3 5		10 5 4
10 5 4		4 4		
8 4 5		10 5		
		2 5		

Formulate hash- and sorting-based algorithms in pseudocode which implement the above mentioned operators. Use pseudocode syntax based on the sample code found on slides 15 and 23 in slide set 12. You are also allowed to use operations like `sort(<table>)` or `[not] in <hashtable>`. If you use unusual operators in your pseudocode, please include a short note on their semantics.

Note: Disregard *blocked I/O* in your pseudocode implementation.

2. [15 Points] **Join Operators in PostgreSQL**

Nested Loop Join: Disable Hash- and Merge Join to answer the following questions!

```
set enable_hashjoin=off;
set enable_mergejoin=off;
```

- (a) Create tables `one` and `many` as provided in `one-many.sql`. As long as there are no keys or indexes created on both tables, a query to join them using *Nested Loop Join* will not terminate in reasonable time. Use `EXPLAIN` to show the plan of the following query Q :

```
SELECT *
FROM   one AS o, many AS m
WHERE  o.a = m.a
```

- Based on the estimated rows, how often would the *Join Filter* `o.a = m.a` be evaluated?

- (b) A PRIMARY KEY index on `one(a)` supports the *Nested Loop Join* on query Q . How?

```
ALTER TABLE one ADD CONSTRAINT one_a PRIMARY KEY (a);
ALTER TABLE many ADD FOREIGN KEY (a) REFERENCES one(a);
ANALYZE one; ANALYZE many;
```

- Show the plan using `EXPLAIN ANALYZE` and explain briefly.
- (c) An additional `PRIMARY KEY` index on `many(a,c)` again improves the query performance.
- ```
ALTER TABLE many ADD CONSTRAINT many_a_c PRIMARY KEY (a,c);
ANALYZE one; ANALYZE many;
```
- Why is only one of both indexes used, while the other table is accessed using a `Seq Scan`?
  - Why is table `many` with index `many_a_c` (and not table `one` with `one_a`) preferred as inner join table here?
- (d) How does the following modification of query `Q` benefit from both indexes, instead?
- ```
SELECT *
FROM   one AS o, many AS m
WHERE  o.a = m.a
ORDER BY m.a
```

Hash Join: Re-enable Hash- and Merge Join to answer the following questions!

```
set enable_hashjoin=on;
set enable_mergejoin=on;
```

- (e) If available, a *Hash Join* is used to answer the equi-join query `Q`. Show the plan using `EXPLAIN (VERBOSE, ANALYZE, BUFFERS)`.
- Why is table `one` (and not table `many`) chosen as the build table here?
 - Why are the indexes `one_a` and `many_a_c` not used here?
- (f) *Hash Join* builds a temporary hash table and thus suffers when `work_mem` is reduced. `set work_mem='64kB'` (instead of default `'4MB'`) and re-execute the query. The *Hash Join* performance decreases significantly. Use the output of `EXPLAIN (VERBOSE, ANALYZE, BUFFERS)` to compare it with `2e` and explain why.
- (g) Since a part of the hash table is stored in-memory even for low `work_mem`, the actual performance of `2f` can highly depend on the data distribution of the probed table. Table `many_skewed` in `one-many.sql` provides a variant of table `many` with a heavily skewed distribution on column `a`. Examine columns `n_distinct`, `most_common_vals` and `most_common_freqs` in table `pg_stats`¹ to show statistics about the distribution of values for both, attribute `a` in `many` and `a` in `many_skewed`.
- Give a short comparison.
- Execute the following query on `work_mem='64kB'` and compare its plan to `2f`.²
- The I/O on temporary tables is reduced. To which number and why?

```
SELECT *
FROM   one AS o, many_skewed AS m
WHERE  o.a = m.a
```

Merge Join: Disable Hash Join to answer the following question!

```
set enable_hashjoin=off;
```

- (h) When we enforce *Merge Join* in `2f`, we can observe that *Merge Join* performs quite better than the *Hash Join*.
- How does the *Merge Join* make use of indexes `one_a` and `many_a_c`?
 - Why is memory no crucial factor here, so that *Merge Join* can outperform the *Hash Join* on `work_mem='64kB'`?

¹<https://www.postgresql.org/docs/current/static/view-pg-stats.html>

²You may have to execute each query twice to avoid inconsistencies in caching.