

DB 2

03 – Wide Table Storage

Summer 2018

Torsten Grust
Universität Tübingen, Germany

1 : Q_2 — Querying a Wider Table

The next **SQL probe** Q_2 looks just Q_1 . We query a wider table now, however:

```
SELECT t.*           -- *  $\equiv$  access all columns of row t
FROM ternary AS t
```

Retrieve all rows (in some arbitrary order) and all columns of table **ternary**, a three-column table created by SQL DDL statement

```
CREATE TABLE ternary (a int NOT NULL,
                       b text NOT NULL, -- variable width
                       c float);         -- may be NULL
```

```
db2=# DROP TABLE IF EXISTS ternary;
db2=# CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c float);

db2=# INSERT INTO ternary(a, b, c)
      SELECT i,
             md5(i::text),
             log(i)
      FROM   generate_series(1, 1000, 1) AS i;

-- Q2: Retrieve all rows (in arbitrary order) and all columns of table ternary
db2=# SELECT t.*
      FROM   ternary AS t;

-- Equivalent to Q2
db2=# TABLE ternary;

db2=# EXPLAIN VERBOSE
      SELECT t.*
      FROM ternary AS t;
```

Using **EXPLAIN** on Q_2

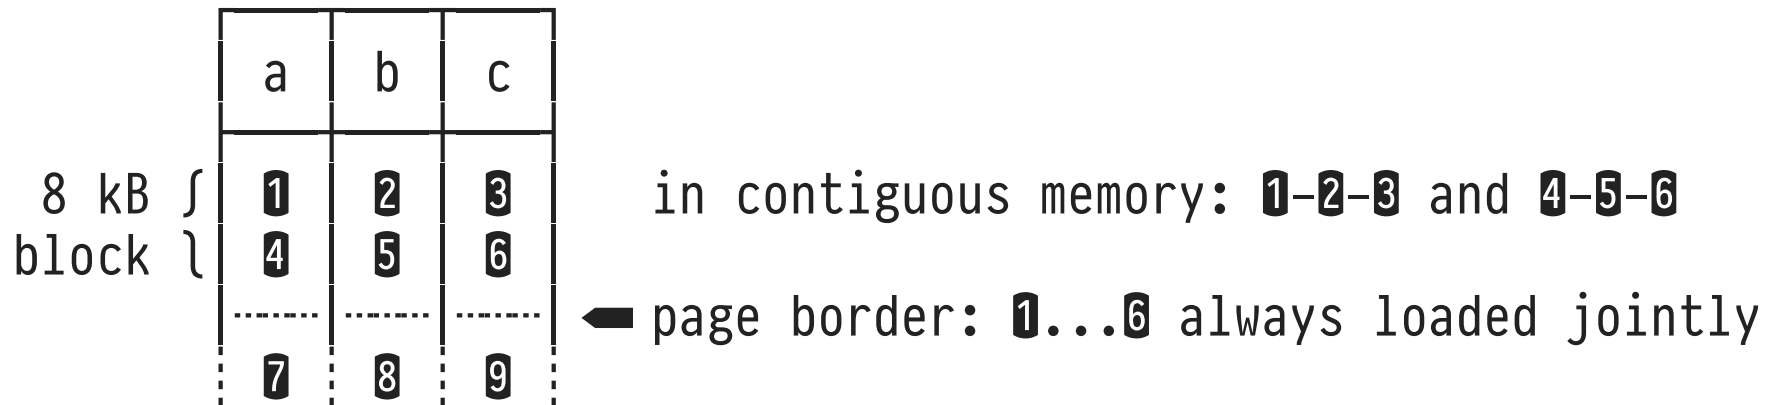
```
EXPLAIN VERBOSE
SELECT t.*
FROM   ternary AS t;
```

QUERY PLAN
Seq Scan on public.ternary t (cost=0.00..20.00 rows=1000 width=45) Output: a, b, c ◀

- Each row **t** carries multiple columns (**a**, **b**, **c**).
- **Seq Scan** scans wider rows now, *average* width: 45 bytes = 4 (int) + 33 (text) + 8 (float) bytes.
 - Column **b** of type text leads to **variable-width rows** in general.

PostgreSQL: Row Storage

- PostgreSQL implements **row storage**: all columns of a row *t* are held in contiguous memory (\equiv same heap file page):

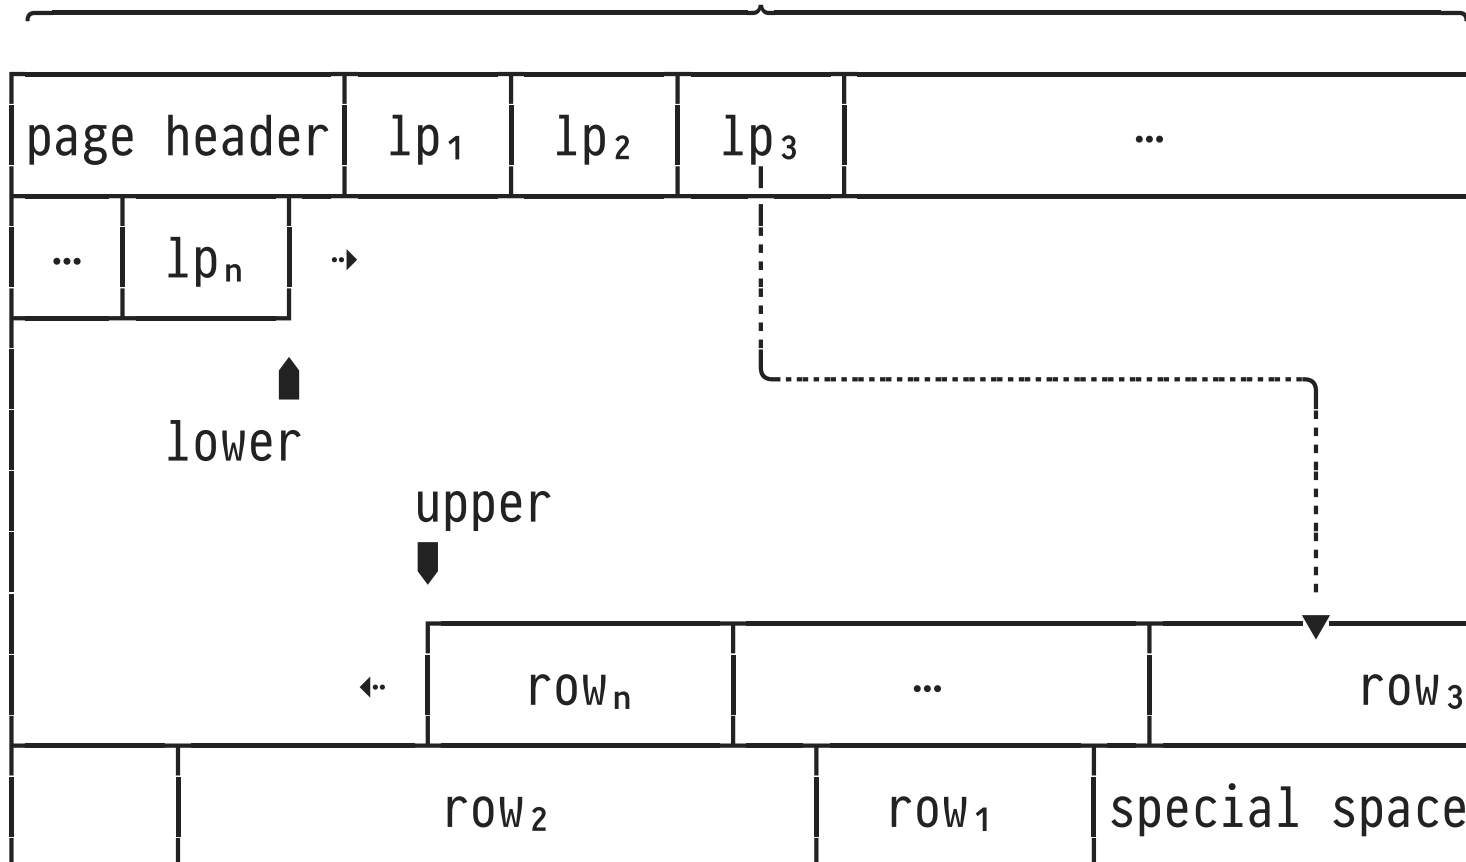


- Loading one heap file page reads **all columns** of all contained rows (recall: block I/O), regardless of whether the query uses *t.** or *t.a* to access the row.

The Innards of a Heap File Page



8 kB block



- lp_i : row pointer
- →/←: direction of growth
- upper - lower: free page space
- row_i : payload

The Innards of a Heap File Page

Comments on the previous slide:

- **Page header** (24 bytes) carries page meta information.
- **Special space** is unused on regular table heap file pages (but used on index pages → later).
- Page is full if pointers **lower** and **upper** meet (row pointers and payload grow towards $\cdots\rightarrow\leftarrow\cdots$ each other).
- **Row pointer** (or: line pointer, 4 byte) lp_i points to row_i , admits **variable-width rows** and **intra-page row relocation** (→ row updates).
- Internal structure of row payloads row_i addressed later.

The Innards of a Heap File Page

PostgreSQL comes with extension `pageinspect` that provides an “X-ray for heap file pages”:

```
CREATE EXTENSION IF NOT EXISTS pageinspect;

-- inspect page header (first 24 bytes)
SELECT *
FROM    page_header(get_raw_page('ternary', <page>));

-- inspect row pointers (lp_i)
SELECT *
FROM    heap_page_items(get_raw_page('ternary', <page>));
```


1 Inspect page header:

```
db2=# CREATE EXTENSION IF NOT EXISTS pageinspect;
```

```
db2=# SELECT t.ctid, t.* FROM ternary AS t;
```

ctid	a	b	c
(0,1)	1	c4ca4238a0b923820dcc509a6f75849b	0
(0,2)	2	c81e728d9d4c2f636f067f89cc14862c	0.301029995663981
(0,3)	3	eccbc87e4b5ce2fe28308fd9f2a7baf3	0.477121254719662
(0,4)	4	a87ff679a2f3e71d9181a67b7542122c	0.602059991327962
(0,5)	5	e4da3b7fbbce2345d7772b0674a318d5	0.698970004336019
[...]			
(9,34)	997	ec5aa0b7846082a2415f0902f0da88f2	2.99869515831166
(9,35)	998	9ab0d88431732957a618d4a469a0d4c3	2.99913054128737
(9,36)	999	b706835de79a2b4e80506f582af3676a	2.99956548822598
(9,37)	1000	a9b7ba70783b617e9998dc4dd82eb3c5	3

← 37 rows on page 9

```
db2=# SELECT * FROM page_header(get_raw_page('ternary', 0));
```

lsn	checksum	flags	lower	upper	special	pagesize	version	prune_xid
1/D4073838	0	0	452	488	8192	8192	4	0

488 - 452 = 36 bytes of free space
(36 < 45 bytes row size: cannot fit more rows)

special @ page end
(⇒ no special space)

```
db2=# SELECT * FROM page_header(get_raw_page('ternary', 9));
```

lsn	checksum	flags	lower	upper	special	pagesize	version	prune_xid
1/D408A420	0	0	172	5528	8192	8192	4	0

5528 - 172 = 5356 bytes of free space

(172 - 24) bytes / 4 bytes = 37 row pointers ≡ 37 rows on page 9

page header per row pointer

Cross-checking with the free space information for table [ternary](#):

```
db2=# VACUUM ternary;
db2=# SELECT * FROM pg_freespace('ternary');
```

blkno	avail	
0	32	└─ approximates 36 free bytes in 32-byte units
1	32	
2	32	
3	32	
4	32	
5	32	
6	32	
7	32	
8	32	
9	5344	└─ approximates 5356 free bytes in 32-byte units

2 Inspect row pointers (on page 0):

```
db2=# SELECT lp, lp_off, lp_len, t_hoff, t_ctid, t_infomask :: bit(16), t_infomask2
FROM heap_page_items(get_raw_page('ternary', 0));
```

lp	lp_off	lp_len	t_hoff	t_ctid	t_infomask	t_infomask2
1	8120	72	24	(0,1)	0000100100000010	3
2	8048	72	24	(0,2)	0000100100000010	3
3	7976	72	24	(0,3)	0000100100000010	3
[...]						
106	560	72	24	(0,106)	0000100100000010	3
107	488	72	24	(0,107)	0000100100000010	3

Also see <https://www.postgresql.org/docs/10/static/storage-page-layout.html>.

- **lp**: number i of lp_i
- **lp_off**: location/offset of row_i on the page
- **lp_len**: width of tuple (includes row header of 24 bytes): 24 bytes header + 45 bytes payload = 69 bytes \approx 72 bytes (rounded to 8 bytes)
- **t_hoff**: offset to row payload data (beyond header and NULL bitmap)
- **t_ctid**: row ID of this row or its next newer version (MVCC)
 - versions of the same row form a chain connected by their **t_ctid** fields, chain ends when **t_ctid** points to the current row
- **t_infomask**:
 - `xxxxxxxxxxxx0`: does tuple have any NULL attribute values?
 - `xxxxxxxxxxxx0x`: does tuple have any variable-length attributes?
 - `xxxxxxxxxxxx0xx`: have attributes of this tuple been stored externally?
 - `xx0xxxxxxxxxxxx`: is this an UPDATeD version of the tuple?
- **t_infomask2**: number of attributes (lower 11 bits + additional flags)

3 Cross-check in on-disk heap file for contents of the third row (lp = 3):

```
db2=# SELECT oid FROM pg_database WHERE datname = 'db2';
```

oid
71857

```
db2=# SELECT relfilenode FROM pg_class WHERE relname = 'ternary';
```

relfilenode
80247

```
db2=# show data_directory;
```

data_directory
/Users/grust/Library/Application Support/Postgres/var-10

```
$ cd '/Users/grust/Library/Application Support/Postgres/var-10/base/71857'
```

```
$ hexdump -C -n 72 -s 7976 80247
```

00001f28	58 31 00 00 00 00 00 00	00 00 00 00 00 00 00 00	X1.....
00001f38	03 00 03 00 02 09 18 00	03 00 00 00 43 65 63 63Cecc
00001f48	62 63 38 37 65 34 62 35	63 65 32 66 65 32 38 33	bc87e4b5ce2fe283
00001f58	30 38 66 64 39 66 32 61	37 62 61 66 33 00 00 00	08fd9f2a7baf3...
00001f68	fd d5 4f 96 27 89 de 3f		..0.'...?
00001f70			

```
db2=# SELECT t.ctid, t.* FROM ternary AS t LIMIT 3;
```

ctid	a	b	c
(0,1)	1	c4ca4238a0b923820dcc509a6f75849b	0
(0,2)	2	c81e728d9d4c2f636f067f89cc14862c	0.301029995663981
(0,3)	3	eccbc87e4b5ce2fe28308fd9f2a7baf3	0.477121254719662

2 : Q_2 — Querying a Wider Table



Recall SQL probe Q_2 :

```
SELECT t.*           -- *  $\equiv$  access all columns of row t
FROM   ternay AS t
```

It is expected that the retrieval of all columns via $t.*$ has consequences for a column-oriented DBMS. We need to **touch and synchronize multiple column vectors**.

Create and populate table `ternary` on MonetDB:

```
$ mclient -d scratch
Welcome to mclient, the MonetDB/SQL interactive terminal (Jul2017-SP4)
Database: MonetDB v11.27.13 (Jul2017-SP4), 'scratch'
Type \q to quit, \? for a list of available commands
auto commit mode: on
sql>DROP TABLE IF EXISTS ternary;
sql>CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c double);
```

```
sql>INSERT INTO ternary(a, b, c)
      SELECT value, md5(value), log(value)
      FROM   generate_series(1, 1001);
```

```
sql>SELECT t.* FROM ternary AS t LIMIT 5;
```


a	b	c
1	c4ca4238a0b923820dcc509a6f75849b	0
2	c81e728d9d4c2f636f067f89cc14862c	0.6931471805599453
3	eccbc87e4b5ce2fe28308fd9f2a7baf3	1.0986122886681098
4	a87ff679a2f3e71d9181a67b7542122c	1.3862943611198906
5	e4da3b7fbbce2345d7772b0674a318d5	1.6094379124341003

Using **EXPLAIN** on Q_2

MAL program for Q_2 , shortened and formatted:

```
⋮
X_4      := sql.mvc();
C_5 :bat[:oid] := sql.tid(X_4, "sys", "ternary");
X_25:bat[:dbl] := sql.bind(X_4, "sys", "ternary", "c",...);
X_31:bat[:dbl] := algebra.projection(C_5, X_25);
X_18:bat[:str] := sql.bind(X_4, "sys", "ternary", "b",...);
X_24:bat[:str] := algebra.projection(C_5, X_18);
X_8  :bat[:int] := sql.bind(X_4, "sys", "ternary", "a",...);
X_17:bat[:int] := algebra.projection(C_5, X_8);
⋮
<create schema of result table>
⋮
sql.resultSet(..., X_17, X_24, X_31);
```

N-ary vs Decomposed Storage Model (NSM/DSM)



MonetDB follows the **Decomposed Storage Model (DSM)** and represents n -ary tables using **full vertical fragmentation**:

a	b	c
a_1	b_1	c_1
a_2	b_2	c_2
a_3	b_3	c_3
a_4	b_4	c_4
a_5	b_5	c_5

NSM (n -ary table)

a:bat[: τ_1]		b:bat[: τ_2]		c:bat[: τ_3]		
head	tail	head	tail	head	tail	
0@0	a_1	0@0	b_1	0@0	c_1	
1@0	a_2	1@0	b_2	1@0	c_2	
2@0	a_3	2@0	b_3	2@0	c_3	← "row" 2@0
3@0	a_4	3@0	b_4	3@0	c_4	
4@0	a_5	4@0	b_5	4@0	c_5	

DSM (n binary tables)

Turn Heads 90°

≡ row
storage

a	b	c
1	2	3
4	5	6
7	8	9

≡ column
storage

a	b	c
1	4	7
2	5	8
3	6	9

in contiguous memory:

1-2-3, 4-5-6, 7-8-9

- Both types of DBMS exhibit strengths/weaknesses for different classes of workloads (→ OLTP vs. OLAP).

Positional BAT Joins

Reconstruction of the n -ary table requires n BATs that are **synchronized on their heads** (\equiv identical cardinality).

- Conceptually: $(n-1)$ -fold natural \bowtie on the head columns.
- Implemented: synchronized scan of the n tail columns:

head	tail		head	tail		head	tail	
0@0	a_1	←	0@0	b_1	←	0@0	c_1	←
1@0	a_2	↓	1@0	b_2	↓	1@0	c_2	↓
2@0	a_3		2@0	b_3		2@0	c_3	

- See variadic MAL builtin `io.print(...,...)`, for example.

EXPLAIN Q_2 in `mclient` SQL REPL, then replay plan at `mserver5` console:

```
sql>EXPLAIN SELECT t.* FROM ternary AS t;
```

```
| mal
```

```
function user.s16_1():void;
  X_1:void := querylog.define("explain select t.* from ternary as t;", "default_pipe", 41:int);
  X_33 := bat.new(nil:str);
  X_39 := bat.new(nil:int);
  X_37 := bat.new(nil:int);
  X_36 := bat.new(nil:str);
  X_35 := bat.new(nil:str);
  X_4 := sql.mvc();
  C_5:bat[:oid] := sql.tid(X_4, "sys", "ternary");
  X_25:bat[:dbl] := sql.bind(X_4, "sys", "ternary", "c", 0:int);
  X_31 := algebra.projection(C_5, X_25);
  X_18:bat[:str] := sql.bind(X_4, "sys", "ternary", "b", 0:int);
  X_24 := algebra.projection(C_5, X_18);
  X_8:bat[:int] := sql.bind(X_4, "sys", "ternary", "a", 0:int);
  X_17 := algebra.projection(C_5, X_8);
  X_40 := bat.append(X_33, "sys.t");
  X_42 := bat.append(X_35, "a");
  X_44 := bat.append(X_36, "int");
  X_46 := bat.append(X_37, 32:int);
  X_48 := bat.append(X_39, 0:int);
  X_50 := bat.append(X_40, "sys.t");
  X_51 := bat.append(X_42, "b");
  X_53 := bat.append(X_44, "clob");
  X_55 := bat.append(X_46, 0:int);
  X_56 := bat.append(X_48, 0:int);
  X_57 := bat.append(X_50, "sys.t");
  X_58 := bat.append(X_51, "c");
  X_60 := bat.append(X_53, "double");
  X_62 := bat.append(X_55, 53:int);
  X_64 := bat.append(X_56, 0:int);
  sql.resultSet(X_57, X_58, X_60, X_62, X_64, X_17, X_24, X_31);
end user.s16_1;
[...]
```

Replay plan at `mserver5` console:

```
sql.init();
sql      := sql.mvc();

ternary:bat[:oid] := sql.tid(sql, "sys", "ternary");
c0      :bat[:dbl] := sql.bind(sql, "sys", "ternary", "c", 0:int);
c       := algebra.projection(ternary, c0);
b0      :bat[:str] := sql.bind(sql, "sys", "ternary", "b", 0:int);
b       := algebra.projection(ternary, b0);
a0      :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
a       := algebra.projection(ternary, a0);

io.print(a,b,c)  ◀ prints three-column table (+ leading void column)
```

`sql.resultSet(...,a,b,c)` creates ternary result table (of name `sys.t`, see row variable `t`) and adds column name and type information.