

# DB 2

---

04 – Row Internals

Summer 2018

Torsten Grust  
Universität Tübingen, Germany

## 1 | $Q_3$ — Projecting on Columns

---

**SQL probe**  $Q_3$  projects on selected columns only (column **b** of table **ternary** is “projected away”):

```
SELECT t.a, t.c      -- access some columns of row t
FROM   ternary AS t
```

Retrieve all rows. Unpack/navigate the row and **extract selected columns**. Recall table **ternary**:

```
CREATE TABLE ternary (a int NOT NULL,
                       b text NOT NULL, -- variable width
                       c float);         -- may be NULL
```

Recreate and populate table `ternary`. Set some `c` values to `NULL`, though:

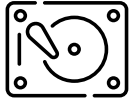
```
DROP TABLE IF EXISTS ternary;
CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c float);

INSERT INTO ternary(a, b, c)
  SELECT i,
         md5(i::text),
         CASE WHEN i % 10 = 0 THEN NULL ELSE log(i) END
  FROM   generate_series(1, 1000, 1) AS i;

-- Q3: Retrieve all rows (in arbitrary order) and selected columns
SELECT t.a, t.c
FROM   unary AS t;

EXPLAIN VERBOSE
SELECT t.a, t.c
FROM   ternary AS t;
```

## Using **EXPLAIN** on $Q_3$



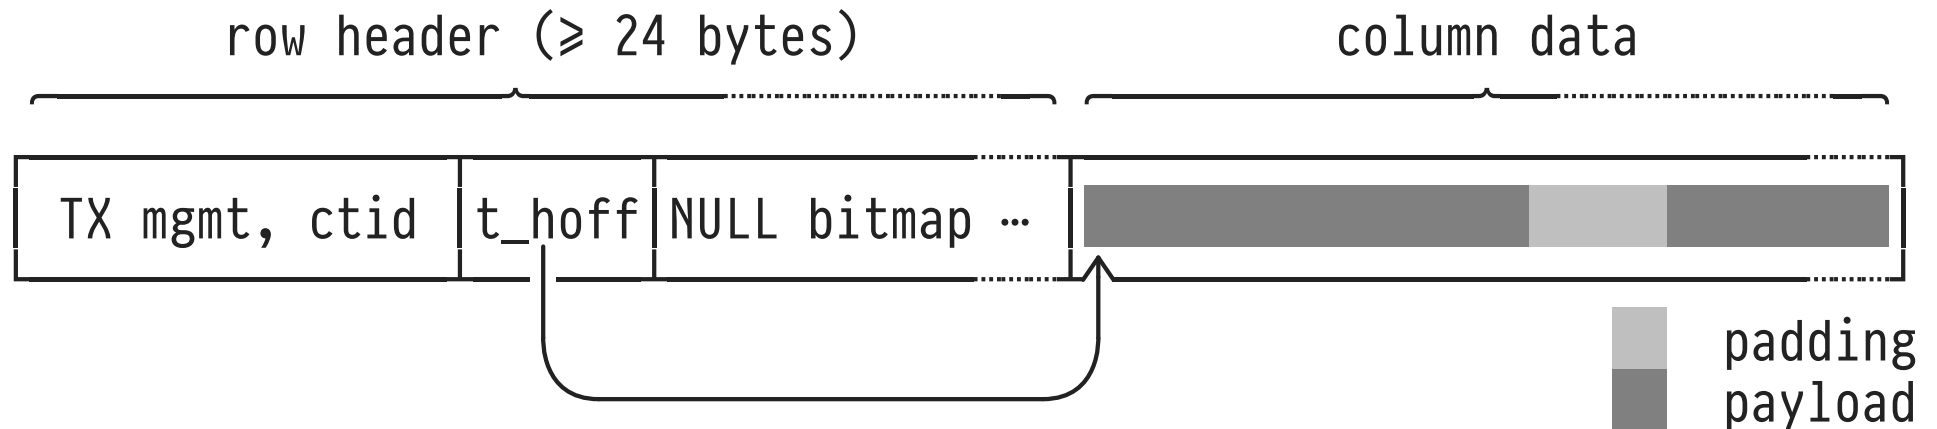
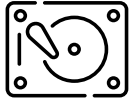
```
EXPLAIN VERBOSE
SELECT t.a, t.c
FROM   ternary AS t;
```

### QUERY PLAN

Seq Scan on public.ternary t (cost=0.00..20.00 rows=1000 width=12) Output: a, c
--


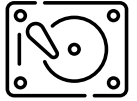
- For each row  $t$ , only columns  $a$  and  $c$  are extracted.
- **Seq Scan** emits narrower rows now, *average* width: 12 bytes = 4 (int) + 8 (float) bytes.
- Estimated cost of 20.00 unchanged from  $Q_2$ :  $Q_3$  does *not* scan fewer data pages ( $\rightarrow$  row storage).

## 2 : Internal Layout of a PostgreSQL Row



- NULL bitmap is of variable length (1 bit per column), offset **t\_hoff** points to first byte of row payload data.
- **NB:** **EXPLAIN**'s **width=w** reports payload bytes only.

# Padding and Alignment



- CPU and memory subsystem require **alignment**: value of width  $n$  bytes is stored at address  $a$  with  $a \bmod n = 0$ .<sup>1</sup>
- $\Rightarrow$  Pad payload such that each column starts at properly aligned offset (PostgreSQL: see table [pg\\_attribute](#)).

<sup>1</sup> Non-aligned data access incur performance penalties (multiple accesses) or even exceptions.



Check table `pg_attribute` for alignment requirements of the column type of table `ternary`:

```
SELECT a.attnum, a.attname, a.attlen, a.attstorage, a.attalign
FROM   pg_attribute AS a
WHERE  a.attrelid = 'ternary'::regclass AND a.attnum > 0
ORDER BY a.attnum;
```


attnum	attname	attlen	attstorage	attalign
1	a	4	p	i
2	b	-1	x	i
3	c	8	p	d

◀ align like an integer (word) = 4 byte

◀ align like a double-word = 8 byte

column width (-1: variable)     p: stored plain  
                                   x: inline (may be compressed) or external (TOAST)

## “Column Tetris”



Padding may lead to substantial space overhead. If viable, reorder columns to tightly pack rows and avoid padding:

```
CREATE TABLE padded (  
  d int2,  
  a int8,  
  e int2,  
  b int8,  
  f int2,  
  c int8)
```

48

```
CREATE TABLE packed (  
  a int8  -- int8: 8-byte aligned  
  b int8  
  c int8  
  d int2  -- int2: 2-byte aligned  
  e int2  
  f int2)
```

30 (+2) column data width

- **+2**: Rows start at **MAXALIGN** offsets ( $\equiv 8$  on 64-bit CPUs).



Demonstrate effect of "column tetris":

```
DROP TABLE IF EXISTS padded;
DROP TABLE IF EXISTS packed;
CREATE TABLE padded (d int2, a int8, e int2, b int8, f int2, c int8);
CREATE TABLE packed (a int8, b int8, c int8, d int2, e int2, f int2);

INSERT INTO padded(d,a,e,b,f,c) SELECT 0,i,0,i,0,i FROM generate_series(1,1000000) AS i;
INSERT INTO packed(a,b,c,d,e,f) SELECT i,i,i,0,0,0 FROM generate_series(1,1000000) AS i;
```

```
SELECT COUNT(*) FROM pg_freespace('padded');
```

count
9346

```
SELECT COUNT(*) FROM pg_freespace('packed');
```

count
7353

← table 'packed' uses fewer pages (79%)

```
SELECT lp, lp_off, lp_len, t_hoff, t_ctid FROM heap_page_items(get_raw_page('padded',0));
```

lp	lp_off	lp_len	t_hoff	t_ctid
1	8120	72	24	(0,1)
2	8048	72	24	(0,2)
[...]				
106	560	72	24	(0,106)
107	488	72	24	(0,107)

← 107 rows per page, each row has 24 + 48 = 72 bytes

```
SELECT lp, lp_off, lp_len, t_hoff, t_ctid FROM heap_page_items(get_raw_page('packed',0));
```

lp	lp_off	lp_len	t_hoff	t_ctid
1	8136	54	24	(0,1)
2	8080	54	24	(0,2)
[...]				
135	632	54	24	(0,135)
136	576	54	24	(0,136)

← 136 rows per page, each row uses 24 + 30 = 54 bytes  
(MAXALIGN: a row starts every 56 bytes on the page)

**EXPLAIN VERBOSE SELECT p.\* FROM padded AS p;**

QUERY PLAN

Seq Scan on public.padded p (cost=0.00..19346.00 rows=1000000 width=30) ← lie! diregards padding  
Output: d, a, e, b, f, c

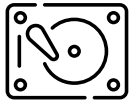
**EXPLAIN VERBOSE SELECT p.\* FROM packed AS p;**

QUERY PLAN

Seq Scan on public.packed p (cost=0.00..17353.00 rows=1000000 width=30) ← lie! diregards padding  
Output: a, b, c, d, e, f



DBMS expects to **perform less work on table 'packed'**

# NULL (Non-)Storage



a	b	c
1	abc	0.1
⋮	⋮	⋮
2	def	NULL
⋮	⋮	⋮

any column NULL?



NULL bitmap ( $\lceil \text{table width} / 8 \rceil$  bytes)

- **NULL** values are represented by **0** bits in a **NULL bitmap** (bitmap is present only if the row indeed contains a **NULL**).

Use extension `pageinspect` to check `NULL`-related row details:

```
SELECT lp, lp_off, lp_len, t_hoff, t_ctid, t_infomask::bit(1) AS "any NULL?", t_bits
FROM heap_page_items(get_raw_page('ternary',0));
```

lp	lp_off	lp_len	t_hoff	t_ctid	any NULL?	t_bits
1	8120	72	24	(0,1)	0	□
2	8048	72	24	(0,2)	0	□
3	7976	72	24	(0,3)	0	□
4	7904	72	24	(0,4)	0	□
5	7832	72	24	(0,5)	0	□
6	7760	72	24	(0,6)	0	□
7	7688	72	24	(0,7)	0	□
8	7616	72	24	(0,8)	0	□
9	7544	72	24	(0,9)	0	□
10	7480	61	24	(0,10)	1	11000000

[...]

column a starts at offset 24 in row (NULL bitmap = MAXALIGN padding = 1 byte)

61 = 24 + 4 + 33 (no intra-row padding)

header    a    b

          int    text (md5 hash is 32 chars long, 1 byte: string length)

Use table `padded` to insert and then inspect an all-NULL row (what will length/NULL bitmap look like)?


```
INSERT INTO padded(a,b,c,d,e,f) VALUES (NULL, NULL, NULL, NULL, NULL, NULL);

SELECT p.ctid FROM padded AS p WHERE (a,b,c,d,e,f) IS NULL; -- ≡ a IS NULL AND b IS NULL AND ...
```

ctid
(9345,86)


```
SELECT lp, lp_off, lp_len, t_hoff, t_ctid, t_infomask::bit(1) AS "any NULL?", t_bits
FROM   heap_page_items(get_raw_page('padded',9345))
WHERE  lp = 86;
```

lp	lp_off	lp_len	t_hoff	t_ctid	any NULL?	t_bits
86	2048	24	24	(9345,86)	1	00000000



## Column Access (Projection)

---

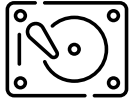


- If  $t$  denotes a row, **column access** — denoted using dot notation  $t.a$  — is the most common operation in SQL query expressions.
  - A typical SQL query will perform multiple column accesses per row (in **SELECT**, **WHERE**, **GROUP BY**, ... clauses), potentially millions of times during evaluation of a single query.
- Even tiny savings in processing effort (here: CPU time) will add up and can lead to substantial benefits.<sup>2</sup>

<sup>2</sup> This is a recurring theme in DBMS implementation. The larger the table cardinalities, the more worthwhile “micro optimizations” become.

## Column Access (Projection)


---



- PostgreSQL: access  $i$ th column of a row using C routine `slot_getattr(i)`:
  1. Has value for column  $i$  been cached? If so, immediately return value.
  2. Check bit for  $i$ th column in NULL bitmap (if present): if 0, immediately return `NULL`.
  3. Scan row payload data **from left to right** for all columns  $k \leq i$ :
    - Use type of column  $k$  to decode payload bytes.
    - Skip over contents if column  $k$  has variable width.
    - **Cache decoded value** for column  $k$  for subsequent `slot_getattr(k)` calls.

## Column Access: PostgreSQL's `slot_getattr()`

---



See PostgreSQL source code (a prime example of readable, consistent, well-documented C code — go read it!):

- File `src/backend/access/common/heaptuple.c`:

```
Datum
slot_getattr(TupleTableSlot *slot, int attnum, bool *isnull)
{
    /* step 1. check cache for column attnum ( $\equiv i$ ) */
    /* step 2. check NULL bitmap */
    :
    /*
     * Extract the attribute, along with any preceding attributes.
     */
    slot_deform_tuple(slot, attnum);
    :
}
```

- `slot_deform_tuple()` does the hard decoding work (step 3.)



PostgreSQL source, file [src/backend/access/common/heaptuple.c](#):

```
/*
 * slot_getattr
 *   This function fetches an attribute of the slot's current tuple.
 *   It is functionally equivalent to heap_getattr, but fetches of
 *   multiple attributes of the same tuple will be optimized better,
 *   because we avoid O(N^2) behavior from multiple calls of
 *   nocachegetattr(), even when attcacheoff isn't usable.
 */
Datum
slot_getattr(TupleTableSlot *slot, int attnum, bool *isnull)
{
    HeapTuple tuple = slot->tts_tuple;
    TupleDesc tupleDesc = slot->tts_tupleDescriptor;
    HeapTupleHeader tup;

    /*
     * system attributes are handled by heap_getsysattr
     */
    if (attnum <= 0)
    {
        if (tuple == NULL) /* internal error */
            elog(ERROR, "cannot extract system attribute from virtual tuple");
        if (tuple == &(slot->tts_minhdr)) /* internal error */
            elog(ERROR, "cannot extract system attribute from minimal tuple");
        return heap_getsysattr(tuple, attnum, tupleDesc, isnull);
    }

    /*
     * fast path if desired attribute already cached
     */
    if (attnum <= slot->tts_nvalid)
    {
        *isnull = slot->tts_isnull[attnum - 1];
        return slot->tts_values[attnum - 1];
    }

    /*
     * return NULL if attnum is out of range according to the tupdesc
     */
    if (attnum > tupleDesc->natts)
    {
```

```

    *isnull = true;
    return (Datum) 0;
}

/*
 * otherwise we had better have a physical tuple (tts_nvalid should equal
 * natts in all virtual-tuple cases)
 */
if (tuple == NULL)      /* internal error */
    elog(ERROR, "cannot extract attribute from empty tuple slot");

/*
 * return NULL if attnum is out of range according to the tuple
 *
 * (We have to check this separately because of various inheritance and
 * table-alteration scenarios: the tuple could be either longer or shorter
 * than the tupdesc.)
 */
tup = tuple->t_data;
if (attnum > HeapTupleHeaderGetNatts(tup))
{
    *isnull = true;
    return (Datum) 0;
}

/*
 * check if target attribute is null: no point in groveling through tuple
 */
if (HeapTupleHasNulls(tuple) && att_isnull(attnum - 1, tup->t_bits))
{
    *isnull = true;
    return (Datum) 0;
}

/*
 * If the attribute's column has been dropped, we force a NULL result.
 * This case should not happen in normal use, but it could happen if we
 * are executing a plan cached before the column was dropped.
 */
if (TupleDescAttr(tupleDesc, attnum - 1)->attisdropped)
{
    *isnull = true;
    return (Datum) 0;
}

/*
 * Extract the attribute, along with any preceding attributes.

```

```

    */
    slot_deform_tuple(slot, attnum);

    /*
     * The result is acquired from tts_values array.
     */
    *isnull = slot->tts_isnull[attnum - 1];
    return slot->tts_values[attnum - 1];
}

```

Experiment: What's the contribution of `slot_deform_tuple()` to the overall PostgreSQL runtime, when the DBMS evaluates a query?

Prepare a large variant of the `ternary` table:

```

DROP TABLE IF EXISTS ternary_10M;
CREATE TABLE ternary_10M(a int NOT NULL, b text NOT NULL, c float);

INSERT INTO ternary_10M(a, b, c)
SELECT i,
       md5(i::text),
       CASE WHEN i % 10 = 0 THEN NULL ELSE log(i) END
FROM   generate_series(1, 10000000, 1) AS i;

```

Use `Activity Monitor` to find out the PID of the `postgres` process that is used to perform work when a query is submitted (query takes about 3 seconds):

```
EXPLAIN (VERBOSE, ANALYZE) SELECT t.b, t.c FROM ternary_10M AS t;
```

Use macOS' `sample` utility to capture the call stack of `postgres` during a 5 second interval while the above query is executing:

```

(in other terminal) $ sample postgres 5
You have access to multiple processes named postgres:
  a) 56806 /Applications/Postgres.app/Contents/Versions/10/bin/postgres
  [...]
  h) 57167 /Applications/Postgres.app/Contents/Versions/10/bin/postgres
Which process? (letter or PID) <h>

```

```

(in psql -d db2)
EXPLAIN (VERBOSE, ANALYZE) SELECT t.b, t.c FROM ternary_10M AS t;

```


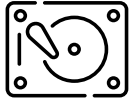
On other terminal/shell find output like this:

Sort by top of stack, same collapsed (when  $\geq 5$ ):

poll (in libsystem_kernel.dylib)	5894	
slot_deform_tuple (in postgres)	394	◀ A top contributor
__commpage_gettimeofday (in libsystem_kernel.dylib)	370	
mach_absolute_time (in libsystem_kernel.dylib)	350	
read (in libsystem_kernel.dylib)	232	
ExecInterpExpr (in postgres)	156	
heapgettup_pagemode (in postgres)	118	
slot_getsomeattrs (in postgres)	83	
HeapTupleSatisfiesMVCC (in postgres)	73	
ExecProject (in postgres)	68	

[...]

## Alternative Layout of Row Payload: Fixed-Width First



- **Separate fixed- from variable-width** payload data at `⋮`:
  - `■■⋮` : fixed-width columns `a`, `c` (types `int`, `double`) + fixed-width pointers `b*` to variable-width columns
  - `⋮■■` : variable-width value for column `b` (type `text`)
- $\Rightarrow$  Can calculate offsets of fixed-width columns **at query compile time**, no left-to-right scanning at run time.

#### Representation of `NULL`?

- Presence of `NULL` left of `:` may spoil static precalculation of column offsets. Option: represent `NULL` using special bit pattern that is as wide as a regular value (if column access performance more important than space savings).

Row navigation and column access is costly. Experiment:

```
DROP TABLE IF EXISTS ternary_10M;  
CREATE TABLE ternary_10M (a int NOT NULL, b text NOT NULL, c float);  
  
INSERT INTO ternary_10M(a, b, c)  
  SELECT i,  
         md5(i::text),  
         CASE WHEN i % 10 = 0 THEN NULL ELSE log(i) END  
  FROM   generate_series(1, 10000000, 1) AS i;
```

Perform two queries that retrieve and deliver the same column contents:

- Query 1 accesses columns `c`, `b`, `a` individually.
- Query 2 reproduces rows as is (fast-path, no actual access to individual columns).

```
-- 1  
EXPLAIN (VERBOSE, ANALYZE)  
  SELECT t.c, t.b, t.a  
  FROM ternary_10M AS t;
```

#### QUERY PLAN

```
Seq Scan on public.ternary_10m t (cost=0.00..192593.44 rows=10000044 width=45) (actual time=0.030..2306.229 rows=10000000 loops=1)  
  Output: b, c, a  
  Planning time: 0.033 ms  
  Execution time: 2827.920 ms ◀
```

```
-- 2  
EXPLAIN (VERBOSE, ANALYZE)  
  SELECT t.*           -- also OK: t.a, t.b, t.c ≡ t.*  
  FROM ternary_10M AS t;
```

#### QUERY PLAN

```
Seq Scan on public.ternary_10m t (cost=0.00..192593.44 rows=10000044 width=45) (actual time=0.066..1328.628 rows=10000000 loops=1)  
  Output: a, b, c  
  Planning time: 0.037 ms  
  Execution time: 1839.703 ms ◀
```

### 3 : $Q_3$ — Projecting on Columns

---



Column **b** of table `ternary(a,b,c)` is irrelevant for the projection query  $Q_3$ :

```
SELECT t.a, t.c      -- access some columns of row t
FROM   ternary AS t
```

We expect the column-oriented DBMS to **exclusively touch the relevant columns**. The wider the input table (and the less columns are accessed), the higher the expected benefit over the row-based DBMS.




Create and populate table `ternary` on MonetDB:

```
$ mclient -d scratch
Welcome to mclient, the MonetDB/SQL interactive terminal (Jul2017-SP4)
Database: MonetDB v11.27.13 (Jul2017-SP4), 'scratch'
Type \q to quit, \? for a list of available commands
auto commit mode: on

DROP TABLE IF EXISTS ternary;
CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c double);


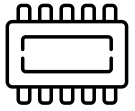
INSERT INTO ternary(a, b, c)
  SELECT value, md5(value), log(value)
  FROM   generate_series(1, 1001);
```



MAL program for  $Q_3$ , shortened and formatted (compare with the MAL program for  $Q_2$ ):

```
⋮
X_4      := sql.mvc();
C_5 :bat[:oid] := sql.tid(X_4, "sys", "ternary");
X_18:bat[:dbl] := sql.bind(X_4, "sys", "ternary", "c", ...);
X_24:bat[:dbl] := algebra.projection(C_5, X_18);
X_8 :bat[:int] := sql.bind(X_4, "sys", "ternary", "a", ...);
X_17:bat[:int] := algebra.projection(C_5, X_8);
⋮
<create schema of result table>
⋮
sql.resultSet(..., X_17, X_24);
```

# Don't Need it? Don't Load it!



- **100%** of the data loaded by the column-based DBMS is useful for query evaluation.