

DB 2

05 – Row Updates

Summer 2018

Torsten Grust
Universität Tübingen, Germany

1 | Q_4 — Row Update

SQL probe Q_4 uses SQL DML statements (**INSERT**, **UPDATE**, **DELETE**) to alter the state of a table:

INSERT INTO ternary	UPDATE ternary	DELETE FROM ternary
SELECT ...	SET $c = e_1$	WHERE $a = e_2$
FROM ...	WHERE $a = e_2$	

INSERT: evaluate query to construct new rows.

UPDATE, **DELETE**: query the table and identify the affected rows.

Modify table storage to reflect the row updates.¹

¹ We still assume that the table has *no* associated index structures.

(Re-)Create table `ternary`:

```
DROP TABLE IF EXISTS ternary;  
CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c float);
```

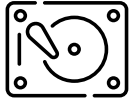
```
INSERT INTO ternary(a, b, c)  
  SELECT i,  
         md5(i::text),  
         log(i)  
FROM    generate_series(1, 1000, 1) AS i;
```

-- Q4: Perform row insertions/updates/deletions (⚠ no ANALYZE here, don't alter the table now):

```
EXPLAIN VERBOSE  
INSERT INTO ternary(a,b,c)  
  SELECT t.a, 'Han Solo', t.c  
FROM    ternary AS t;
```

```
EXPLAIN VERBOSE  
UPDATE ternary AS t  
  SET   c = -1  
 WHERE t.a = 982;
```

Using **EXPLAIN** on Q_4 : **INSERT**



EXPLAIN VERBOSE

```
INSERT INTO ternary(a,b,c)
  SELECT t.a, 'Han Solo', t.c
FROM   ternary AS t;
```

QUERY PLAN

```
Insert on public.ternary (cost=0.00..20.00 rows=1000 width=44)
-> Seq Scan on public.ternary t (cost=0.00..20.00 rows=1000 width=44)
    Output: t.a, 'Han Solo'::text, t.c
```

- **Seq Scan** scans table **ternary** to construct rows to be inserted, feeds 1000 rows into **Insert** for insertion.
- Width of inserted rows (over-)estimated to be 44 bytes = 4 (int) + 32 (text) + 8 (float) bytes.

Insert rows into table `ternary` (⚠ uses `EXPLAIN` without `ANALYZE` to not actually modify the table at this point):

EXPLAIN VERBOSE

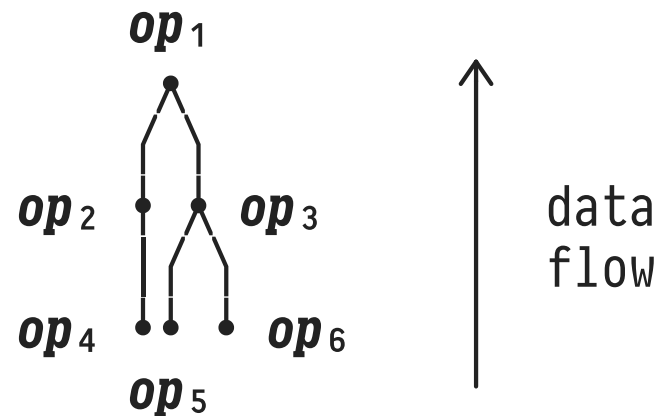
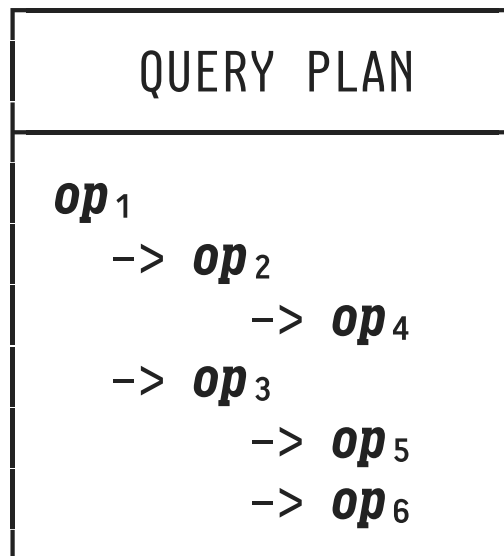
```
INSERT INTO ternary(a,b,c)
SELECT t.a, 'Han Solo', t.c
FROM   ternary AS t;
```

QUERY PLAN
Insert on public.ternary (cost=0.00..20.00 rows=1000 width=44) -> Seq Scan on public.ternary t (cost=0.00..20.00 rows=1000 width=44) Output: t.a, 'Han Solo'::text, t.c

Any column of type `text` – regardless of length – appears to account for 32 bytes in the `width` output of `EXPLAIN`. Actual rows on page have the expected width.

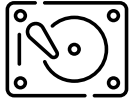
Reading Complex **EXPLAIN** Outputs

- **EXPLAIN** uses symbol `->` and indentation to visualize larger, **tree-shaped** query evaluation plans:



- Read plans “inside out”, root ***op*₁** delivers query result.

Using **EXPLAIN** on Q_4 : **UPDATE**



```
EXPLAIN VERBOSE
UPDATE ternary AS t
  SET    c = -1
  WHERE t.a = 982;
```

QUERY PLAN

```
Update on public.ternary t (cost=0.00..22.50 rows=1 width=51)
-> Seq Scan on public.ternary t (cost=0.00..22.50 rows=1 width=51)
    └─> Output: a, b, '-1'::double precision, ctid
        Filter: (t.a = 982)
```

- **Seq Scan** emits complete rows (only **c** was updated).
- Additionally feeds row ID (**ctid**) into **Update** to identify the affected row(s).

Update and delete rows in table `ternary` (⚠ uses `EXPLAIN` without `ANALYZE` to not actually modify the table at this point):

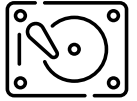
EXPLAIN VERBOSE

```
UPDATE ternary AS t
SET    c = -1
WHERE  t.a = 982;
```

QUERY PLAN
Update on public.ternary t (cost =0.00..22.50 rows =1 width=51) -> Seq Scan on public.ternary t (cost =0.00..22.50 rows =1 width=51) Output: a, b, '-1':: <u>double</u> precision, ctid Filter: (t.a = 982)

Width of 51 > 45 accounts for the additional 6 bytes of the emitted `ctid` (row ID) field.

Using **EXPLAIN** on Q_4 : **DELETE**



```
EXPLAIN VERBOSE
```

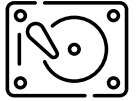
```
DELETE FROM ternary AS t
WHERE t.a = 982;
```

QUERY PLAN

```
Delete on public.ternary t (cost=0.00..22.50 rows=1 width=6)
-> Seq Scan on public.ternary t (cost=0.00..22.50 rows=1 width=6)
    Output: ctid ◀
    Filter: (t.a = 982)
```

- **Seq Scan** returns affected row IDs (of 6 bytes each) only.
- We turn to **Filter** (makes scan skip non-qualifying rows) later in this course.

2 : How do Row Updates Alter the Table Storage?



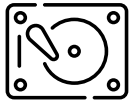
Let us take a closer look at how plan operator **Update** alters the target table's heap file pages. We find:²

- Rows are *not* updated in-place. A **new version of the row is created** — original and updated row co-exist.
- Any database user (query, application) sees exactly one version of any row at any time. Different users may see different row versions.
- A separate **VACUUM** (“garbage collection”) step collects and removes old versions that cannot be seen by any user.

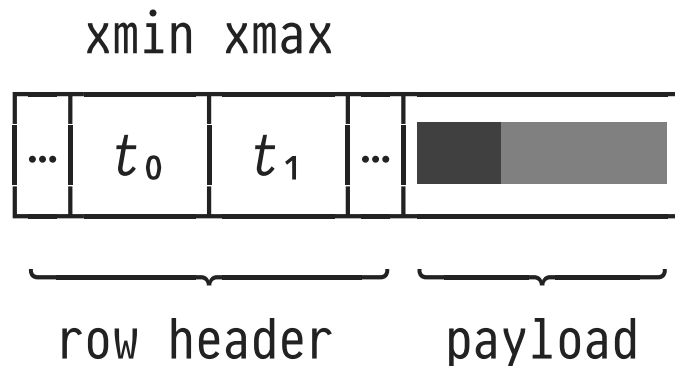
² This implementation of **Update** is typical for all DBMS that implement **Multi-Version Concurrency Control (MVCC)**. We discuss MVCC later in this course.



Row Visibility and Timestamps



1. Each row carries two **timestamps** — `xmin` and `xmax` — that mark its first and last time of existence.
2. Each query/update is executed at some timestamp `T` which defines the rows that are **visible** to the operation:



row is visible for any operation
with timestamp $t_0 \leq T < t_1$

($t_1 = \infty$: row has not been updated yet)

- DBMS uses system-wide virtual timestamps (*transaction IDs*), see PostgreSQL built-in function `txid_current()`.

Explore row version chaining and row visibility for table `ternary`. Table has 9 pages, where the last page has space to spare. Update a row on page 9:

```
-- recreate a clean ternary table
DROP TABLE IF EXISTS ternary;
CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c float);
INSERT INTO ternary(a, b, c)
  SELECT i,
         md5(i::text),
         log(i)
  FROM   generate_series(1, 1000, 1) AS i;

-- ❶ list rows on page 9
SELECT t.ctid, t.* FROM ternary AS t WHERE t.ctid >= '(9,1)';
```

ctid	a	b	c
(9,1)	964	8065d07da4a77621450aa84fee5656d9	2.98407703390283
(9,2)	965	eeb69a3cb92300456b6a5f4162093851	2.98452731334379
[...]			
(9,18)	981	287e03db1d99e0ec2edb90d079e142f3	2.99166900737995
(9,19)	982	fec8d47d412bcbeece3d9128ae855a7a	2.99211148778695
(9,20)	983	6aab1270668d8cac7cef2566a1c5f569	2.99255351783214
[...]			
(9,36)	999	b706835de79a2b4e80506f582af3676a	2.99956548822598
(9,37)	1000	a9b7ba70783b617e9998dc4dd82eb3c5	3

← will be updated

```
-- ❷ check row header contents before update
SELECT t_ctid, lp, lp_off, lp_len, t_xmin, t_xmax,
       (t_infomask::bit(16) & b'0010000000000000')::int::bool AS "updated row?",
       (t_infomask2::bit(16) & b'0100000000000000')::int::bool AS "has been HOT updated?"
FROM   heap_page_items(get_raw_page('ternary', 9));
```

t_ctid	lp	lp_off	lp_len	t_xmin	t_xmax	updated row?	has been HOT updated?
(9,1)	1	8120	72	12708	0	f	f
(9,2)	2	8048	72	12708	0	f	f
[...]							
(9,18)	18	6896	72	12708	0	f	f
(9,19)	19	6824	72	12708	0	f	f
(9,20)	20	6752	72	12708	0	f	f
[...]							
(9,36)	36	5600	72	12708	0	f	f
(9,37)	37	5528	72	12708	0	f	f

← will be updated

row visible indefinitely, no update yet

```
-- 3 check current transaction ID (= virtual timestamp)
SELECT txid_current();
```

txid_current
12710

 ← current operation timestamp

```
SELECT txid_current();
```

txid_current
12711

 ← current operation timestamp

```
-- 4 update one row
UPDATE ternary AS t SET c = -1 WHERE t.a = 982;
```

```
-- 5 check visible contents of page 9 after update
SELECT t.ctid, t.* FROM ternary AS t WHERE t.ctid >= '(9,1)';
```

ctid	a	b	c
(9,1)	964	8065d07da4a77621450aa84fee5656d9	2.98407703390283
(9,2)	965	eeb69a3cb92300456b6a5f4162093851	2.98452731334379
[...]			
(9,18)	981	287e03db1d99e0ec2edb90d079e142f3	2.99166900737995
(9,20)	983	6aab1270668d8cac7cef2566a1c5f569	2.99255351783214
[...]			
(9,36)	999	b706835de79a2b4e80506f582af3676a	2.99956548822598
(9,37)	1000	a9b7ba70783b617e9998dc4dd82eb3c5	3
(9,38)	982	fec8d47d412bcbeece3d9128ae855a7a	-1

← row (9,19) is gone (≡ invisible)

← updated row

```
-- 6 check row header contents after update
SELECT t_ctid, lp, lp_off, lp_len, t_xmin, t_xmax,
       (t_infomask::bit(16) & b'0010000000000000')::int::bool AS "updated row?",
       (t_infomask2::bit(16) & b'0100000000000000')::int::bool AS "has been HOT updated?"
FROM heap_page_items(get_raw_page('ternary', 9));
```

t_ctid	lp	lp_off	lp_len	t_xmin	t_xmax	updated row?	has been HOT updated?
(9,1)	1	8120	72	12708	0	f	f
(9,2)	2	8048	72	12708	0	f	f
[...]							
(9,18)	18	6896	72	12708	0	f	f
(9,38)	19	6824	72	12708	12713	f	t ←
(9,20)	20	6752	72	12708	0	f	f
[...]							
(9,36)	36	5600	72	12708	0	f	f
(9,37)	37	5528	72	12708	0	f	f
(9,38)	38	5456	72	12713	0	t ←	f

points to itself
(≡ end of chain)

visibility of old row version
is limited

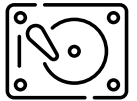
← points to new row version

← updated row

```
db2=# SELECT txid_current();
```

txid_current
12714 ← 12713 < 12714 ⇒ old row version is invisible to us, as expected

Impact of Updates Beyond the Row's Page



- Updates on full pages may lead to row relocation across pages: versions then have row IDs (p_i, lp_i) , (p_{i+1}, lp_{i+1}) where $p_i \neq p_{i+1}$.
 - Traversal of longer update chains may lead to I/O-costly “page hopping.”
 - \Rightarrow Perform **VACUUM** to collect inaccessible old versions. From outside page, point to most recent row directly.
- PostgreSQL optimizes for the good-natured case where $p_i = p_{i+1}$ and indexed row fields have *not* been changed.
 - Such **heap-only tuple (HOT) updates** have page-internal impact only, no maintenance outside page required.

3 | Q_4 — Row Update



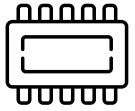
INSERT INTO ternary VALUES (...,...,...)	UPDATE ternary SET c = e_1 WHERE a = e_2	DELETE FROM ternary WHERE a = e_2
---	---	--

UPDATE: affects updated column(s) only.

INSERT, **DELETE**: operate on full rows, all column BATs of table **ternary** will be affected.

- MonetDB uses user-specific **Δ tables** (“delta tables”) to represent changes. Column BATs are *not* modified immediately. Global visibility of changes is delayed.

Using **EXPLAIN** on Q_4 : **DELETE**

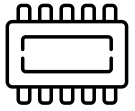


```
sql> EXPLAIN DELETE FROM ternary WHERE a = 981;
```

```
⋮  
ternary      :bat[:oid] := sql.tid(sql, "sys", "ternary");  
a0           :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);  
deleted_rows:bat[:oid] := algebra.thetaselect(a0, ternary, 981:int, "==");  
sql_delete   := sql.delete(sql, "sys", "ternary", deleted_rows);  
⋮
```

- `algebra.thetaselect(b_1, b_2, v, θ)` returns the oids of those rows r in `algebra.projection(b_2, b_1)`, for which predicate `tail(r) θ v` holds.
- `sql.delete(...)` modifies the BAT of currently visible rows (obtainable via `sql.tid(...)`) for table `ternary`.
⚠ However, no column BAT is changed yet.

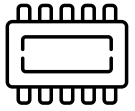
Using **EXPLAIN** on Q_4 : **INSERT**



```
sql> EXPLAIN INSERT INTO ternary(a,b,c) VALUES (1001, 'Han Solo', -2);  
  
:  
sql_append := sql.append(sql, "sys", "ternary", "a", 1001:int);  
sql_append := sql.append(sql, "sys", "ternary", "b", "Han Solo");  
sql_append := sql.append(sql, "sys", "ternary", "c", -2:dbl);  
:
```

- For **ternary(a,b,c)**, a row insert translates into three individual operations on the column BATs.
- **sql.append(...)** saves the inserted value in the **Δ^i table** associated with each column BAT.
 - ⚠ The column BATs do not change yet.

Using **EXPLAIN** on Q_4 : **UPDATE**



```
sql> EXPLAIN UPDATE ternary SET c = -1 WHERE a = 982;
```

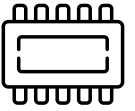
```
⋮
ternary      :bat[:oid] := sql.tid(sql, "sys", "ternary");
a0           :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
updated_rows :bat[:oid] := algebra.thetaselect(a0, ternary, 982:int, "==");
updated_rows_a:bat[:int] := algebra.projection(updated_rows, a0);
updated_c    :bat[:dbl] := algebra.project(updated_rows_a, -1:dbl);
sql_update   := sql.update(sql, "sys", "ternary", "c",
                           updated_rows, updated_c);
⋮
```

- BATs `updated_rows` and `updated_c` contain oids and `c` values of the changed rows.³
- `sql.update(...)` saves these changes in the **Δ^u table** for the BAT of column `c`. ⚠ The column BAT is not changed yet.

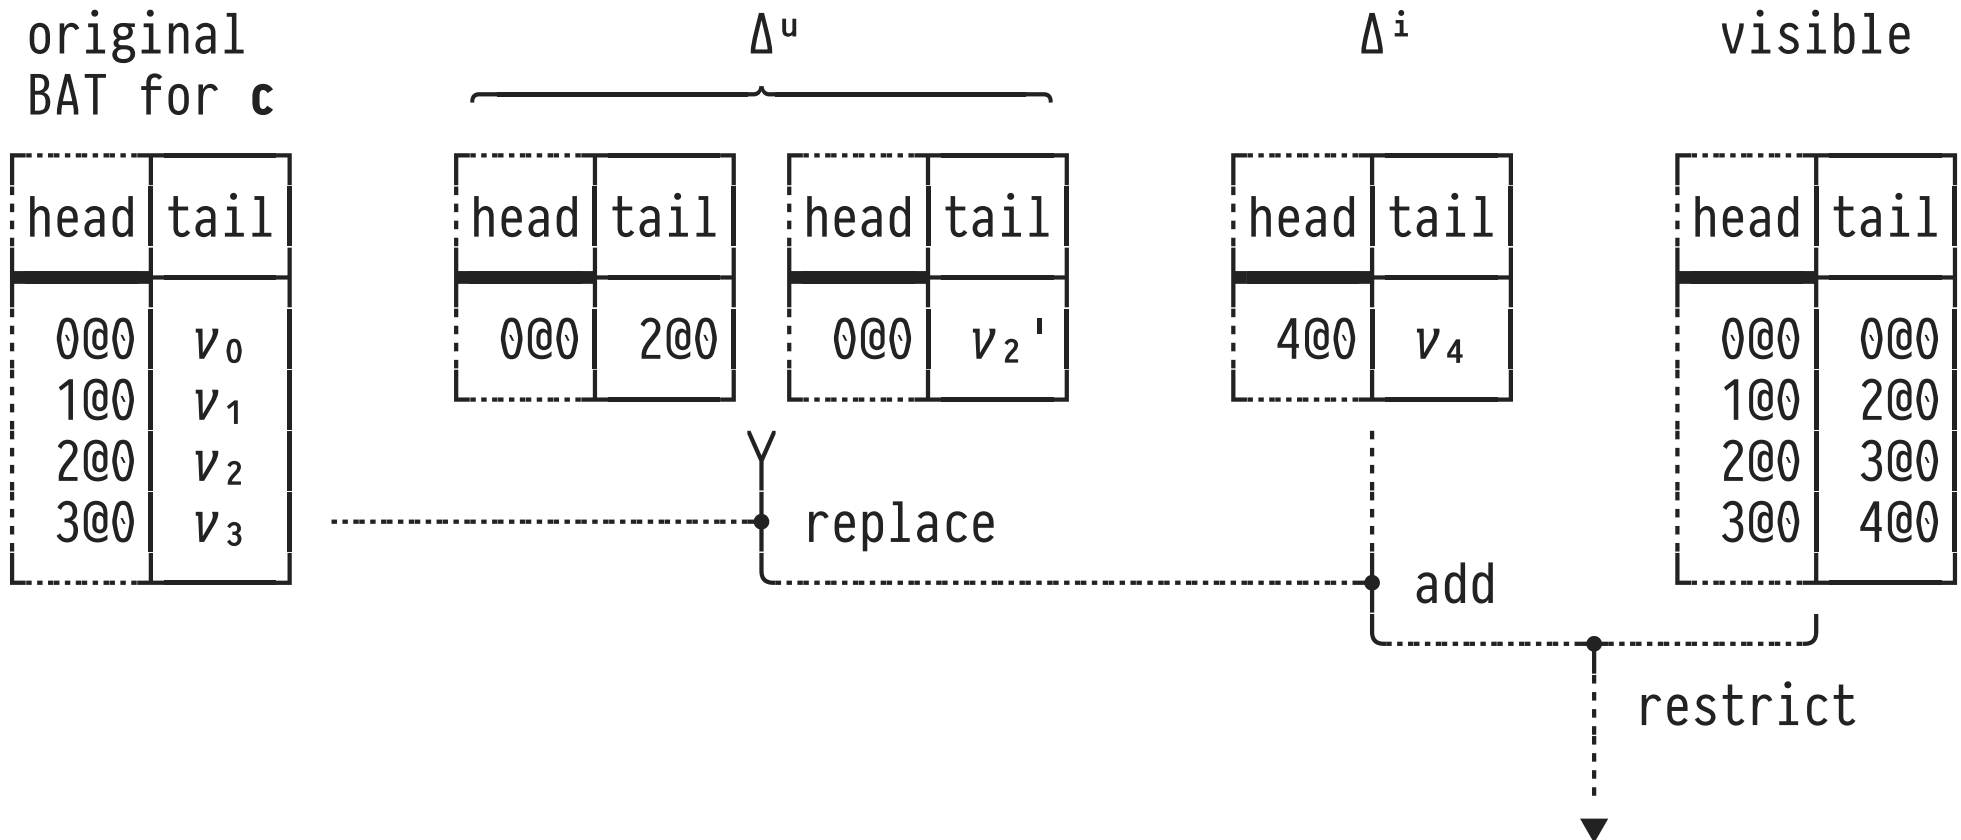
³ `algebra.project(b,v)` returns `b` with all tail values set to `v`.

⚠ The use of BAT `updated_rows_a` above appears unnecessary in this case (could be replaced by `updated_rows`).

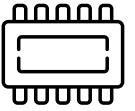
Δ and Visibility Tables



In column c , update v_2 to v_2' , insert v_4 , delete v_1 :

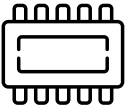


Applying Changes on Demand



- When a query needs to see the changes made to column `c`, apply all changes accumulated in the `c`'s Δ tables:
 1. Load — yet unmodified — column BAT for `c`.
 2. Read `c`'s Δ^u table and perform value replacements (via `bat.replace(...)`).
 3. Read `c`'s Δ^i table and perform value inserts (via `bat.append(...)`).
 4. Restrict `c` to currently visible rows (via `algebra.projection(...)`).
- Make changes permanent only once we want them to be seen globally by all users (`COMMIT` \rightarrow transaction management).

Delay Change Propagation: Disable *AUTO COMMIT*



- To experiment with Δ -based change management, disable MonetDB's default *auto commit* behavior (use `mclient` option `--autocommit` to turn *auto commit* off):

```
DELETE FROM ternary WHERE a = 981;  
INSERT INTO ternary(a,b,c) VALUES (1001, 'Han Solo', -2);  
UPDATE ternary SET c = -1 WHERE a = 982;
```

- Without *auto commit*, changes are still pending at this point. Thus:

```
SELECT t.c           \ reads  $\Delta^u$ ,  $\Delta^i$  (+ visibility table)  
FROM   ternary AS t; \ to reflect changes on t.c
```


Demonstrate Δ tables and delayed change propagation on table `ternary`.

❶ (Re-)Create and populate table `ternary`:

```
$ mclient -d scratch -l sql
```

```
DROP TABLE IF EXISTS ternary;  
CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c double);
```

```
INSERT INTO ternary(a, b, c)  
  SELECT value, md5(value), log(value)  
  FROM   generate_series(1, 1001);
```

```
SELECT t.* FROM ternary AS t LIMIT 5;
```

a	b	c
1	c4ca4238a0b923820dcc509a6f75849b	0
2	c81e728d9d4c2f636f067f89cc14862c	0.6931471805599453
3	eccbc87e4b5ce2fe28308fd9f2a7baf3	1.0986122886681098
4	a87ff679a2f3e71d9181a67b7542122c	1.3862943611198906
5	e4da3b7fbbce2345d7772b0674a318d5	1.6094379124341003

❷ Perform updates on `ternary`, then query column `c`. ⚠ Turn *auto commit* off such that changes accumulate in the Δ tables:

```
$ mclient -d scratch -l sql --autocommit
```

```
sql>EXPLAIN DELETE FROM ternary WHERE a = 981;  
[...]
```

```
  X_9:bat[:int] := sql.bind(X_5, "sys", "ternary", "a", 0:int);  
  C_6:bat[:oid] := sql.tid(X_5, "sys", "ternary");  
  C_21 := algebra.thetaselect(X_9, C_6, 981:int, "==");  
  X_26 := sql.delete(X_5, "sys", "ternary", C_21);  
[...]
```

```
sql>DELETE FROM ternary WHERE a = 981;
```

```
sql>EXPLAIN UPDATE ternary SET c = -1 WHERE a = 982;  
[...]
```

```
  X_10:bat[:int] := sql.bind(X_6, "sys", "ternary", "a", 0:int);  
  C_7:bat[:oid] := sql.tid(X_6, "sys", "ternary");  
  C_22 := algebra.thetaselect(X_10, C_7, 982:int, "==");  
  X_24 := algebra.projection(C_22, X_10);  
  X_28 := algebra.project(X_24, -1:dbl);  
  X_30 := sql.update(X_6, "sys", "ternary", "c", C_22, X_28);  
[...]
```

```
sql>UPDATE ternary SET c = -1 WHERE a = 982;
```

```
sql>EXPLAIN INSERT INTO ternary(a,b,c) VALUES (1001, 'Han Solo', -2);
```

```
[...]
```

```
X_20 := sql.append(X_7, "sys", "ternary", "a", 1001:int);  
X_25 := sql.append(X_20, "sys", "ternary", "b", "Han Solo");  
X_28 := sql.append(X_25, "sys", "ternary", "c", -2:dbl);
```

```
[...]
```

```
sql>INSERT INTO ternary(a,b,c) VALUES (1001, 'Han Solo', -2);
```

```
sql>EXPLAIN SELECT t.c FROM ternary AS t;
```

```
[...]
```

```
C_5:bat[:oid] := sql.tid(X_4, "sys", "ternary");  
X_8:bat[:dbl] := sql.bind(X_4, "sys", "ternary", "c", 0:int);  
(C_13:bat[:oid], X_14:bat[:dbl]) := sql.bind(X_4, "sys", "ternary", "c", 2:int);  
X_11:bat[:dbl] := sql.bind(X_4, "sys", "ternary", "c", 1:int);  
X_16 := sql.delta(X_8, C_13, X_14, X_11);  
X_17 := algebra.projection(C_5, X_16);
```

access Δ^u

access Δ^i

apply Δ^u, Δ^i

restrict to visible

```
[...]
```

```
sql> SELECT t.* FROM ternary AS t;
```

a	b	c
1	c4ca4238a0b923820dcc509a6f75849b	0

```
[...]
```

980	d79aac075930c83c2f1e369a511148fe	6.887552571664617	row with a = 981 deleted
982	fec8d47d412bcbeece3d9128ae855a7a	-1	updated
983	6aab1270668d8cac7cef2566a1c5f569	6.890609120147166	
984	d93ed5b6db83be78efb0d05ae420158e	6.891625897052253	
985	54a367d629152b720749e187b3eaa11b	6.892641641172089	
986	fe7ee8fc1959cc7214fa21c4840dff0a	6.893656354602635	
987	df6d2338b2b8fce1ec2f6dda0a630eb0	6.894670039433482	
988	9908279ebbf1f9b250ba689db6a0222b	6.895682697747868	
989	a1140a3d0df1c81e24ae954d935e8926	6.8966943316227125	
990	4fac9ba115140ac4f1c22da82aa0bc7f	6.897704943128636	
991	692f93be8c7a41525c0baf2076aecfb4	6.898714534329988	
992	860320be12a1c050cd7731794e231bd3	6.899723107284872	
993	7b13b2203029ed80337f27127a9f1d28	6.900730664045173	
994	934815ad542a4a7c5e8a2dfa04fea9f5	6.901737206656574	
995	2bcab9d935d219641434683dd9d18a03	6.902742737158593	
996	0b8aff0438617c055eb55f0ba5d226fa	6.903747257584598	
997	ec5aa0b7846082a2415f0902f0da88f2	6.904750769961838	
998	9ab0d88431732957a618d4a469a0d4c3	6.905753276311464	
999	b706835de79a2b4e80506f582af3676a	6.906754778648554	
1000	a9b7ba70783b617e9998dc4dd82eb3c5	6.907755278982137	
1001	Han Solo	-2	inserted

3 Replay MAL commands on `mclient` console:

Re-create and populate table `ternary`:

```
$ mclient -d scratch -l sql

DROP TABLE IF EXISTS ternary;
CREATE TABLE ternary (a int NOT NULL, b text NOT NULL, c double);

INSERT INTO ternary(a, b, c)
  SELECT value, md5(value), log(value)
  FROM   generate_series(1, 1001);
```

Replay MAL commands:

```
$ mclient -d scratch -l mal

sql.init();
sql := sql.mvc();

# DELETE FROM ternary WHERE a = 981;
ternary      :bat[:oid] := sql.tid(sql, "sys", "ternary");
a0           :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
deleted_rows:bat[:oid] := algebra.thetaselect(a0, ternary, 981:int, "==");
sql_delete   := sql.delete(sql, "sys", "ternary", deleted_rows);

io.print(deleted_rows);
#-----#
# h t # name
# void void # type
#-----#
[ 0@0, 980@0 ]          ◀ oid of the deleted row

# UPDATE ternary SET c = -1 WHERE a = 982;
ternary:bat[:oid] := sql.tid(sql, "sys", "ternary");
io.print(ternary);
#-----#
# h t # name
# void oid # type
#-----#
[...]
[ 978@0, 978@0 ]
[ 979@0, 979@0 ]
[ 980@0, 981@0 ]
[ 981@0, 982@0 ]          ◀ visibility: row with oid 980@0 missing
```

```

[...]
```

```

updated_rows :bat[:oid] := algebra.thetaselect(a0, ternary, 982:int, "==");
updated_rows_a:bat[:int] := algebra.projection(updated_rows, a0);
updated_c     :bat[:dbl] := algebra.project(updated_rows_a, -1:dbl);
sql_update    := sql.update(sql, "sys", "ternary", "c", updated_rows, updated_c);

io.print(updated_rows, updated_c);
#-----#
# t t t # name
# void void dbl # type
#-----#
[ 0@0, 981@0, -1 ]      ◀ Δu: oid of updated row and new c value

# INSERT INTO ternary(a,b,c) VALUES (1001, 'Han Solo', -2);
sql_append := sql.append(sql, "sys", "ternary", "a", 1001:int);
sql_append := sql.append(sql, "sys", "ternary", "b", "Han Solo");
sql_append := sql.append(sql, "sys", "ternary", "c", -2:dbl);

# SELECT t.c FROM ternary AS t;
ternary:bat[:oid] := sql.tid(sql, "sys", "ternary");
c0      :bat[:dbl] := sql.bind(sql, "sys", "ternary", "c", 0:int);

io.print(ternary);
#-----#
# h t # name
# void oid # type
#-----#
[ 0@0, 0@0 ]
[...]
```

[978@0, 978@0]	
[979@0, 979@0]	
[980@0, 981@0]	
[981@0, 982@0]	
[982@0, 983@0]	
[983@0, 984@0]	
[984@0, 985@0]	
[985@0, 986@0]	
[986@0, 987@0]	
[987@0, 988@0]	
[988@0, 989@0]	
[989@0, 990@0]	
[990@0, 991@0]	
[991@0, 992@0]	
[992@0, 993@0]	

```

◀ visibility: row with oid 980@0 still missing
```

```
[ 993@0, 994@0 ]
[ 994@0, 995@0 ]
[ 995@0, 996@0 ]
[ 996@0, 997@0 ]
[ 997@0, 998@0 ]
[ 998@0, 999@0 ]
[ 999@0, 1000@0 ]      ◀ visibility: new row with oid 1000@0 now visible
```

```
io.print(c0);
#-----#
# h t # name
# void dbl # type
#-----#
[ 0@0, 0 ]
[...]
```

[979@0,	6.887552571664617]
[980@0,	6.8885724595653635]
[981@0,	6.889591308354466]
[982@0,	6.890609120147166]
[983@0,	6.891625897052253]
[984@0,	6.892641641172089]
[985@0,	6.893656354602635]
[986@0,	6.894670039433482]
[987@0,	6.895682697747868]
[988@0,	6.8966943316227125]
[989@0,	6.897704943128636]
[990@0,	6.898714534329988]
[991@0,	6.899723107284872]
[992@0,	6.900730664045173]
[993@0,	6.901737206656574]
[994@0,	6.902742737158593]
[995@0,	6.903747257584598]
[996@0,	6.904750769961838]
[997@0,	6.905753276311464]
[998@0,	6.906754778648554]
[999@0,	6.907755278982137]

◀ column BAT for c not updated yet: row 981@0 still visible

◀ column BAT for c not updated yet: row 1000@0 not present yet

```
# access the Δ tables using sql.bind(..., 1) and sql.bind(..., 2)
inserted_c:bat[:dbl]      := sql.bind(sql, "sys", "ternary", "c", 1:int);      ◀ access Δi
(updated_rows:bat[:oid], updated_c:bat[:dbl]) := sql.bind(sql, "sys", "ternary", "c", 2:int); ◀ access Δu

io.print(inserted_c);
#-----#
# h t # name
# void dbl # type
#-----#
[ 1000@0, -2 ]
```

```

io.print(updated_rows, updated_c);
#-----#
# t t t # name
# void oid dbl # type
#-----#
[ 000, 98100, -1 ]

# implement changes (this is equivalent to sql.delta(sql, c0, updated_rows, updated_c, inserted_c))
bat.replace(c0, updated_rows, updated_c, true);    ← apply  $\Delta^u$ 
bat.append(c0, inserted_c);                        ← apply  $\Delta^i$ 
c:bat[:dbl] := algebra.projection(ternary, c0);    ← restrict visibility

io.print(c);
[...]
```

[97800, 6.88653164253051]	
[97900, 6.887552571664617]	
[98000, -1]	← updated
[98100, 6.890609120147166]	
[98200, 6.891625897052253]	
[98300, 6.892641641172089]	
[98400, 6.893656354602635]	
[98500, 6.894670039433482]	
[98600, 6.895682697747868]	
[98700, 6.8966943316227125]	
[98800, 6.897704943128636]	
[98900, 6.898714534329988]	
[99000, 6.899723107284872]	
[99100, 6.900730664045173]	
[99200, 6.901737206656574]	
[99300, 6.902742737158593]	
[99400, 6.903747257584598]	
[99500, 6.904750769961838]	
[99600, 6.905753276311464]	
[99700, 6.906754778648554]	
[99800, 6.907755278982137]	
[99900, -2]	← inserted (also: 1000 rows only since 1 row deleted)