

DB 2

08 – Predicate Evaluation

Summer 2018

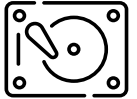
Torsten Grust
Universität Tübingen, Germany

1 | Q₇ — Predicate (or Filter) Evaluation

SQL's **WHERE/HAVING/FILTER** clauses use **expressions of type Boolean (predicates)** to filter rows. Predicates may use Boolean connectives (**AND, OR, NOT**) to build complex filters from simple predicate building blocks:

```
SELECT t.a, t.b
FROM   ternary AS t
WHERE  t.a % 2 = 0 AND [OR] t.c < 1  -- either AND or OR
```

Evaluate predicate for every row **t** scanned. Here:
assume that evaluation of the predicate is *not*
supported by a specific index. (⚠ Index support for
predicates is essential → see upcoming chapters.)



EXPLAIN VERBOSE

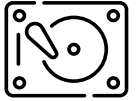
```
SELECT t.a, t.b
FROM   ternary AS t           -- 1000 rows
WHERE  t.a % 2 = 0 AND t.c < 1;
```

QUERY PLAN

Seq Scan on ternary t (cost=... rows=1 ...) (actual time=... rows=5 ...)
Filter: ((c < '1'::double precision) AND ((a % 2) = 0))
Rows Removed by Filter: 995
Planning time: 2.125 ms
Execution time: 1.894 ms

- Filter predicate evaluated during **Seq Scan**.
- Estimated **selectivity** of predicate $1/1000$ (real: $5/1000$).

`t.a % 2 = 0 AND t.c < 1`: An Expression of Type `bool`



- In the absence of index support, use the regular expression interpreter to evaluate predicates:

```
SCAN_FETCHSOME(t, [a, c])
SCAN_VAR(c) ───┐
CONST(1) ───┐
FUNCEXPR_STRICT(<, •, •) ┐
BOOL_AND_STEP_FIRST(•)  # if • = false, immediately yield false
                        #      (∧ semantics: false ∧ p = false)
SCAN_VAR(a) ───┐
CONST(2) ───┐
FUNCEXPR_STRICT(%, •, •) ┐
CONST(0) ───┐
FUNCEXPR_STRICT(=, •, •) ┐
BOOL_AND_STEP_LAST(•)  # yield •      (∧ semantics: true ∧ p = p)
```

- Uses “jumps” in program to implement **Boolean shortcut**.

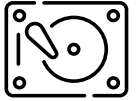
PostgreSQL source code (expression interpreter):


- Boolean shortcut in `BOOL_AND_STEP_FIRST`: `src/backend/executor/execExpInterp.c`, function `ExecInterpExpr()`
- Invocation of expression interpreter in `Seq Scan`: `src/backend/executor/execScan.c`, function `ExecScan()`. See call `ExecQual(qual, econtext)`, preceded by comment:

```
/*
 * check that the current tuple satisfies the qual-clause
 *
 * check for non-null qual here to avoid a function call to ExecQual()
 * when the qual is null ... saves only a few cycles, but they add up
 * ...
 */
if (qual == NULL || ExecQual(qual, econtext))
[...]
```

- Function `ExecQual()` found in `src/include/executor/executor.h`. Can process entire list `ps` of conjunctive clauses:
 - `ExecQual()` handles case `ps = []` and immediately returns **true** (*ex falso quod libet*).

Heuristic Predicate Simplification



- Predicate evaluation effort is multiplied by the number of rows processed. **Even small simplifications add up.**
- PostgreSQL performs basic predicate simplifications:
 - Reduce constant expressions to `true/false`.
 - Apply basic identities (e.g., $\text{NOT}(\text{NOT}(p)) \equiv p$ and $(p \text{ AND } q) \text{ OR } (p \text{ AND } r) \equiv p \text{ AND } (q \text{ OR } r)$).
 - Remove duplicate clauses (e.g., $p \text{ AND } p \equiv p$)
 - Apply De Morgan's laws.
-  These are **heuristics** (expected to improve evaluation time): selectivity is *not yet* taken into account.

Show heuristic predicate simplification at work:

-- ❶ Remove double NOT() + De Morgan:

```
db2=# EXPLAIN VERBOSE
      SELECT t.a, t.b
      FROM   ternary AS t
     WHERE  NOT(NOT(NOT(t.a % 2 = 0 AND t.c < 1)));
```

QUERY PLAN

Seq Scan on public.ternary t (cost=0.00..27.50 rows=1000 width=37) Output: a, b Filter: (((t.a % 2) <> 0) OR (t.c >= '1'::double precision))
--

-- ❷ Inverse distributivity of AND:

```
db2=# EXPLAIN VERBOSE
      SELECT t.a, t.b
      FROM   ternary AS t
     WHERE  (t.a % 2 = 0 AND t.c < 1) OR (t.a % 2 = 0 AND t.c > 2);
```

QUERY PLAN

Seq Scan on public.ternary t (cost=0.00..30.00 rows=5 width=37) Output: a, b Filter: (((t.c < '1'::double precision) OR (t.c > '2'::double precision)) AND ((t.a % 2) = 0))

Machine-Generated Queries and Predicate Simplification

Automatically generated SQL text may differ significantly from human-authored queries. Consider a web search form:

⊗ Search ternary...

a: 42.....

c:

SUBMIT

1. User enters search keys for columns **a** and/or **c**.
2. Web form maps missing keys to **NULL** (interpret as wildcard).
3. DBMS executes parameterized query:

```
SELECT t.*  
FROM   ternary AS t  
WHERE  (t.a = :a OR :a IS NULL)  
       AND (t.c = :c OR :c IS NULL)
```


db2=# EXPLAIN VERBOSE

```
SELECT t.*  
FROM   ternary AS t  
WHERE  (t.a = 42 OR 42 IS NULL)  
AND    (t.c = NULL OR NULL IS NULL);
```

QUERY PLAN

Seq Scan on public.ternary t (cost=0.00..22.50 rows=1 width=45) Output: a, b, c Filter: (t.a = 42)
--


db2=# EXPLAIN VERBOSE

```
SELECT t.*  
FROM   ternary AS t  
WHERE  (t.a = NULL OR NULL IS NULL)  
AND    (t.c = NULL OR NULL IS NULL);
```


QUERY PLAN

Seq Scan on public.ternary t (cost=0.00..20.00 rows=1000 width=45) Output: a, b, c

Heuristics May Not Be Enough



- Heuristics only go so far. The (estimated) **cost** of evaluation may suggest better predicate rewrites:

	(expected) cost
SELECT t.*	
FROM ternary_10m AS t	
WHERE length (btrim(t.b, '0...9')) < length (t.b)	p_1 
OR t.a % 1000 <> 0	p_2 

- With Boolean shortcut it makes a difference which disjunct is evaluated first. (Both predicates not selective, p_1 : 85.9%, p_2 : 99.9% of 10^7 rows pass.)

⇒ Many optimizer decisions indeed **are cost-based**.

Selectivity-based decisions are important in query optimization, but selectivity is secondary for this example. Predicate evaluation cost dominates the overall query evaluation cost:

```
db2=# EXPLAIN ANALYZE
      SELECT *
      FROM ternary_10m AS t
      WHERE length(btrim(t.b, '01234567890')) < length(t.b)
             OR t.a % 1000 <> 0;
```

QUERY PLAN
Seq Scan on ternary_10m (cost=0.00..342594.10 rows=9966711 width=45) (actual time=0.043..14219.828 rows=9998573 loops=1) Filter: ((length(btrim(b, '01234567890'::text)) < length(b)) OR ((a % 1000) <> 0)) Rows Removed by Filter: 1427 Planning time: 0.156 ms Execution time: 14741.087 ms ◀

```
db2=# EXPLAIN ANALYZE
      SELECT *
      FROM   ternary_10m AS t
      WHERE  t.a % 1000 <> 0
             OR length(btrim(t.b, '01234567890')) < length(t.b);
```

QUERY PLAN
Seq Scan on ternary_10m (cost=0.00..342594.10 rows=9966711 width=45) (actual time=0.025..2302.257 rows=9998573 loops=1) Filter: (((a % 1000) <> 0) OR (length(btrim(b, '01234567890'::text)) < length(b))) Rows Removed by Filter: 1427 Planning time: 0.116 ms Execution time: 2803.925 ms ◀

2 : Q_7 — Predicate (or Filter) Evaluation




```
SELECT t.a, t.b
FROM   ternary AS t
WHERE  t.a % 2 = 0 AND [OR] t.c < 1  -- either AND or OR
```

MonetDB can evaluate basic predicates on individual column BATs (here: **a** and **c**) ❶ but then needs to

1. derive the result of composite predicates ❷ and
2. propagate the filter effect to all output columns (here: **a**, **b**) ❸ to form the final selection result.



Using **EXPLAIN** on Q_7 (Boolean Connective: **OR**)




```
sql> EXPLAIN SELECT t.a, t.b
      FROM   ternary AS t
      WHERE  t.a % 2 = 0 OR t.c < 1;

:
ternary :bat[:oid] := sql.tid(sql, "sys", "ternary");
a0      :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
a       :bat[:int] := algebra.projection(ternary, a0);
e1      :bat[:int] := batcalc.%(a, 2:int);           ◀ a % 2
1 p1    :bat[:oid] := algebra.thetaselect(e1, 0:int, "=="); ◀ p1 ≡ a % 2 = 0
c0      :bat[:dbl] := sql.bind(sql, "sys", "ternary", "c", 0:int);
c       :bat[:dbl] := algebra.projection(ternary, c0);
1 p2    :bat[:oid] := algebra.thetaselect(c, 1:dbl, "<");   ◀ p2 ≡ c < 1
2 or     :bat[:oid] := bat.mergeand(p1, p2);             ◀ p1 ∨ p2
b0      :bat[:str] := sql.bind(sql, "sys", "ternary", "b", 0:int);
3 bres   :bat[:str] := algebra.projectionpath(or, ternary, b0); ◀ result col b
3 ares   :bat[:int] := algebra.projection(or, a);         ◀ result col a
:
```

Result of a Predicate \equiv Selection Vectors



Selection Vectors (also: Candidate Lists)



- **Selection vector** `sv`: BAT of type `bat[:oid]`.
 $i@0 \in sv \iff i$ th input row satisfies filter predicate.
- Use `algebra.projection(sv, col)` to propagate filter effect to column `col`.
- Implement Boolean connectives for predicate p_i with `sv_i`:
 - p_1 OR p_2 : `bat.mergeand(sv1, sv2)`
 - p_1 AND p_2 : `algebra.projectionpath(sv2, sv1, •)` with
$$\text{algebra.projectionpath}(sv_2, sv_1, \bullet) \equiv \text{algebra.projection}(sv_2, \text{algebra.projection}(sv_1, \bullet)).$$

Demonstrate evaluation of disjunctive condition in MAL:

```
$ mclient -d scratch -l mal

sql.init();
sql := sql.mvc();
ternary:bat[:oid] := sql.tid(sql, "sys", "ternary");
a0      :bat[:int] := sql.bind(sql, "sys", "ternary", "a", 0:int);
a       :bat[:int] := algebra.projection(ternary, a0);
e1      :bat[:int] := batcalc.%(a, 2:int);  #  $\Leftarrow a \% 2$ 
io.print(e1);
#-----#
# h t # name
# void int # type
#-----#
[ 0@0, 1 ]
[ 1@0, 0 ]
[ 2@0, 1 ]
[ 3@0, 0 ]
[...]
p1      :bat[:oid] := algebra.thetaselect(e1, 0:int, "==");  #  $\Leftarrow p_1 \equiv a \% 2 = 0$ 
io.print(p1);
#-----#
# h t # name
# void oid # type
#-----#
[ 0@0, 1@0 ]
[ 1@0, 3@0 ]
[...]
[ 498@0, 997@0 ]
[ 499@0, 999@0 ]
c0      :bat[:dbl] := sql.bind(sql, "sys", "ternary", "c", 0:int);
c       :bat[:dbl] := algebra.projection(ternary, c0);
p2      :bat[:oid] := algebra.thetaselect(c, 1:dbl, "<");  #  $\Leftarrow p_2 \equiv c < 1$ 
io.print(p2);
#-----#
# h t # name
# void void # type
#-----#
[ 0@0, 0@0 ]       $\Leftarrow c = 0$ 
[ 1@0, 1@0 ]       $\Leftarrow c = 0.6931471805599453$ 

# -----#
# Interlude on range selection
#
#          v <!--≤ hi (false ≡ <)    complement result?
```



```

# algebra.select(col, lo, hi, true, false, false):
#
#           lo </= v (true ≡ ≤)
#
# Return oids of values v in col in the range [lo, hi)

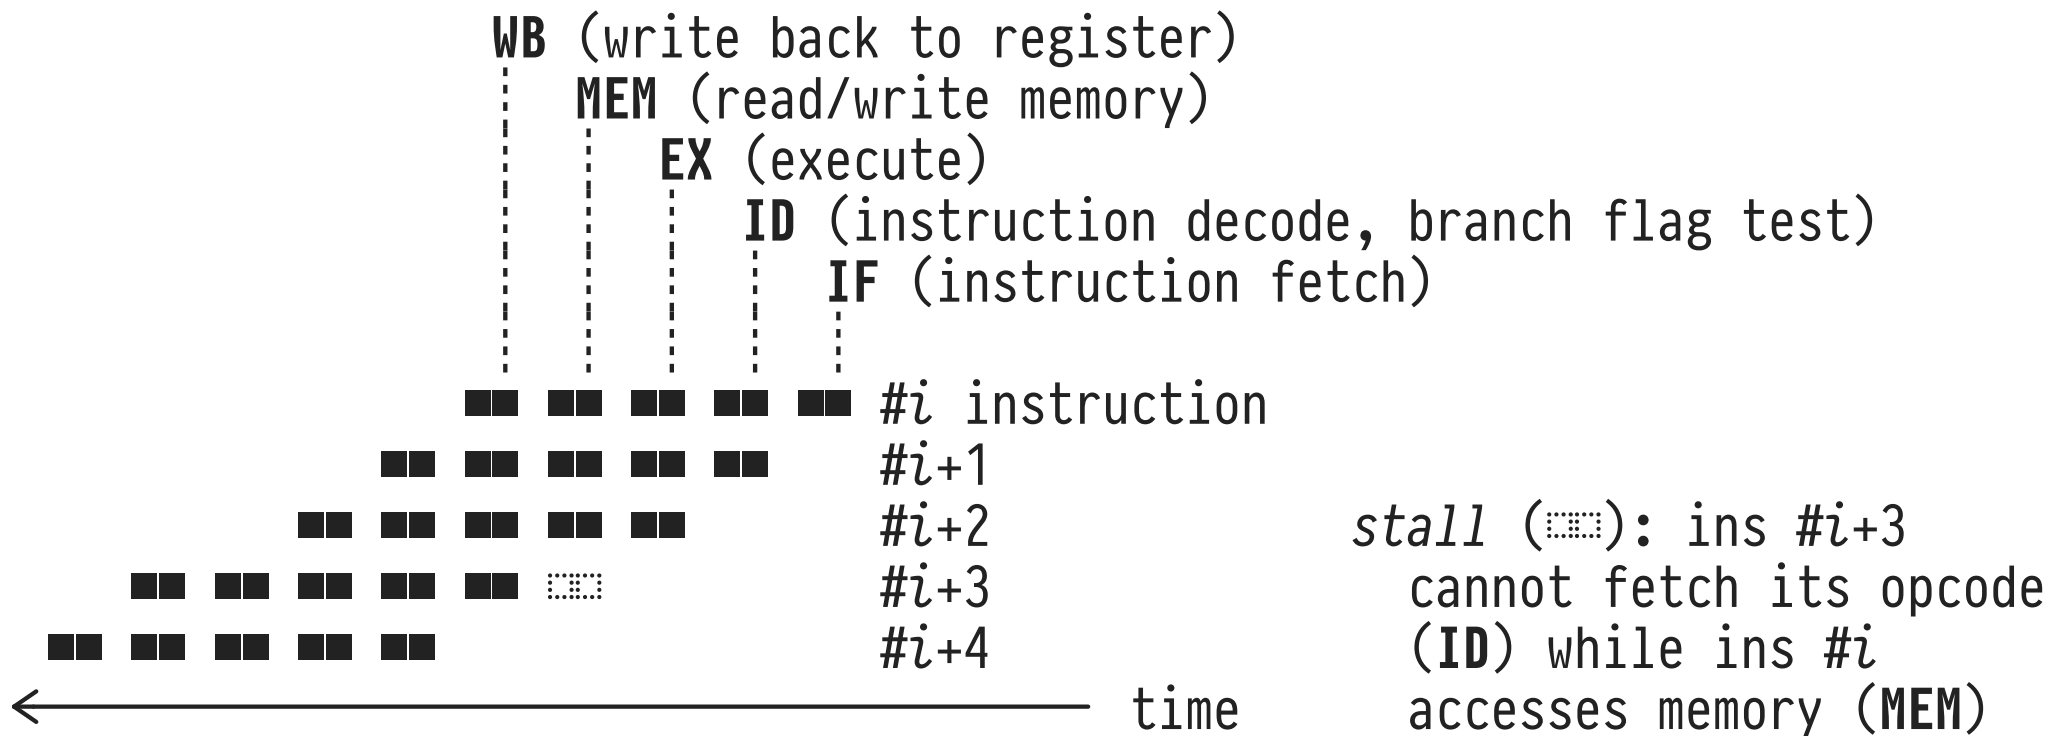
sv := algebra.select(c, 2:dbl, 2.5:dbl, true, false, false); # ◀ 2 ≤ c < 2.5
io.print(sv);
#-----#
# h t # name
# void void # type
#-----#
[ 0@0, 7@0 ]
[ 1@0, 8@0 ]
[ 2@0, 9@0 ]
[ 3@0, 10@0 ]
[ 4@0, 11@0 ]
range := algebra.projection(sv, c);
io.print(range);
# -----#

or := bat.mergeand(p1, p2);      # ◀ p1 ∨ p2
io.print(or);
#-----#
# h t # name
# void oid # type
#-----#
[ 0@0, 0@0 ]   ◀ contributed by e3
[ 1@0, 1@0 ]   ◀ contributed by e2 and e3
[ 2@0, 3@0 ]   ◀ contributed by e2
[... ]
b0      :bat[:str] := sql.bind(sql, "sys", "ternary", "b", 0:int);
bres    :bat[:str] := algebra.projectionpath(or, ternary, b0); # ◀ applies visibility in ternary AND THEN selection vector
ares    :bat[:int] := algebra.projection(or, a);      # ◀ applies selection vector
io.print(ares, bres);
#-----#
# t t t # name
# void int str # type
#-----#
[ 0@0, 1, "c4ca4238a0b923820dcc509a6f75849b" ]
[ 1@0, 2, "c81e728d9d4c2f636f067f89cc14862c" ]
[... ]
[ 499@0, 998, "9ab0d88431732957a618d4a469a0d4c3" ]
[ 500@0, 1000, "a9b7ba70783b617e9998dc4dd82eb3c5" ]

```


Instruction Pipelining in Modern CPUs

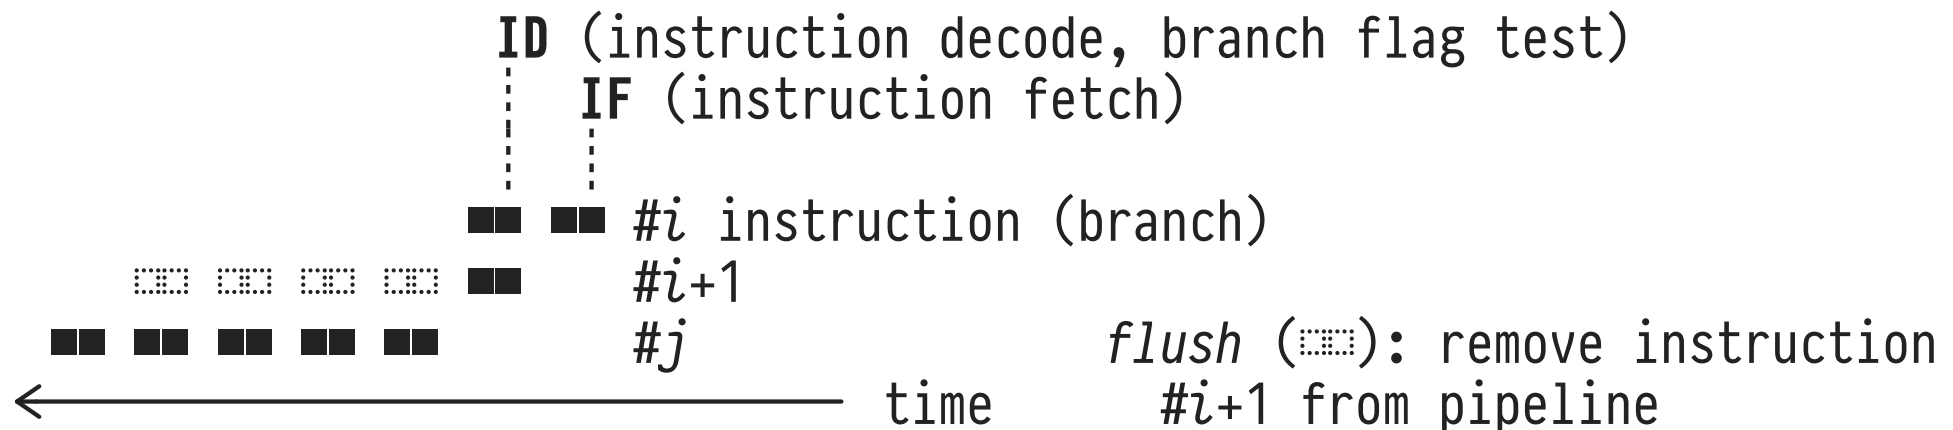
Control flow branches (**for**, but particularly **if**) are a challenge for modern pipelining CPUs:



Branch Taken? Yes, Flush Pipeline

This pipeline decides the outcome of branch $\#i$ (end of **ID**) only *after* instruction $\#i+1$ has already been fetched (**IF**):

- If the branch is taken, **flush** instruction $\#i+1$ from pipeline ☹, instead fetch instruction $\#j$ at jump target:




Branch Prediction: History and Heuristics

CPUs thus try to **predict the outcome of a branch # i** based on **earlier recorded outcomes** of the same branch:

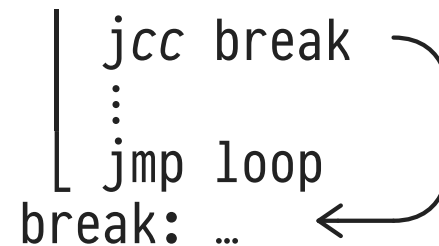
Branch prediction	Fetch instruction
<i>taken</i>	$\#j$
<i>not taken</i>	$\#i+1$

- Also: **heuristics** based on typical control flow patterns:


Predicted *taken*



Predicted *not taken*



Avoiding Branch Mispredictions

- A **mispredicted branch**  leads to
 1. pipeline flushes—effectively a stall—and
 2. (possibly) instruction cache misses.
- The resulting runtime penalty indeed is significant \Rightarrow DBMSs aim to avoid branch mispredictions in tight inner loops:
 - prefer branch-less implementations of query logic,
 - reduce number of random/hard-to-predict branches.

Demonstrate the effects of branch misprediction using a selection `col < v`. Linearly grow `v` from 0...max. Selectivity of predicate `col < v` linearly grows from 0% to 100%.

See file [mat/branch-prediction.c](#), core tight loop:

```
out = 0;
for (int i = 0; i < SIZE; i += 1) {
    if (col[i] < v) {
        sv[out] = i;
        out += 1;
    }
}
```

Compile via (⚠ no `-O` or `-O2` here):

```
$ cc -Wall branch-prediction.c -o branch-prediction
```

QUIZ: Let students speculate about outcome of experiment.

🔧 Sample run:

```
$ ./branch-prediction
0 (selectivity: 0.00%) 33207µs  ◀ predicate always false: branch always taken (perfect to predict, also meets heuristic)
1 (selectivity: 10.00%) 55499µs  } branch outcome increasingly random and thus hard to predict,
2 (selectivity: 20.00%) 82161µs  } branch mispredictions increasingly likely
3 (selectivity: 29.99%) 107224µs }
4 (selectivity: 40.00%) 153250µs }
5 (selectivity: 49.99%) 128811µs } branch outcome essentially arbitrary and thus impossible to predict
6 (selectivity: 59.98%) 120362µs }
7 (selectivity: 69.99%) 103390µs }
8 (selectivity: 80.01%) 83365µs  } branch becomes more and more like to be untaken,
9 (selectivity: 90.01%) 65245µs  } branch mispredictions decrease
10 (selectivity: 100.00%) 48755µs ◀ predicate always true: branch never taken (perfect to predict)
```

🔧 Experiment (1): **Sort** the column vector in ascending order. Column vector is effectively divided into two parts:

1. starting region of column vector where always `col < v`: branch never taken \Rightarrow perfect to predict
2. ending region of column vector where always `col > v`: branch always taken \Rightarrow perfect to predict

(Actual value of `v` immaterial—the larger `v`, the longer `sv` grows which needs a bit of work.)

Sample run for Experiment (1):

```

0 (selectivity: 0.00%) 27578µs  ◀ builds empty selection vector, no memory writes
1 (selectivity: 10.00%) 35127µs
2 (selectivity: 20.00%) 37500µs
3 (selectivity: 29.99%) 40296µs
4 (selectivity: 40.00%) 39552µs
5 (selectivity: 49.99%) 41864µs
6 (selectivity: 59.98%) 43514µs
7 (selectivity: 69.99%) 43268µs
8 (selectivity: 80.01%) 43903µs
9 (selectivity: 90.01%) 46592µs
10 (selectivity: 100.00%) 50677µs  ◀ builds long selection vector, lots of memory writes

```

❸ Experiment (2): A **branch-less** selection.

```

out = 0;
for (int i = 0; i < SIZE; i += 1) {
    sv[out] = i;
    out += (col[i] < v);      ◀ col[i] < v == 1 ⇔ predicate satisfied
}

```

Sample run for Experiment (2):

```

0 (selectivity: 0.00%) 45288µs  } ◀ repeatedly writes to sv[0] only
1 (selectivity: 10.00%) 46625µs
2 (selectivity: 20.00%) 49592µs
3 (selectivity: 29.99%) 47187µs
4 (selectivity: 40.00%) 49891µs  } run time unaffected by predicate selectivity
5 (selectivity: 49.99%) 49738µs
6 (selectivity: 59.98%) 49811µs
7 (selectivity: 69.99%) 46630µs
8 (selectivity: 80.01%) 47391µs
9 (selectivity: 90.01%) 47367µs
10 (selectivity: 100.00%) 52009µs } ◀ writes to entire sv[0...SIZE-1]

```

File `mat/branch-prediction.c`:

```

/* Demonstrate the effects of branch mispredictions for a selection
 * col < v implemented in a tight loop
 */
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <assert.h>

#define MICROSECS(t) (1000000 * (t).tv_sec + (t).tv_usec)

```



```

#define SIZE (16 * 1024 * 1024)
#define STEPS 11

/* comparison of a, b (used in qsort) */
int cmp (const void *a, const void *b)
{
    return *((int*) a) - *((int*) b);
}

int main()
{
    int *col;          /* column vector */
    int *sv;           /* selection vector */

    int out;
    float selectivity;

    struct timeval t0, t1;
    unsigned long duration;

    /* allocate memory */
    col = malloc(SIZE * sizeof(int));
    assert(col);
    sv = malloc(SIZE * sizeof(int));
    assert(sv);

    /* initialize column with (pseudo) random values in interval 0...RAND_MAX */
    srand(42);
    for (int i = 0; i < SIZE; i += 1) {
        col[i] = rand();
    }

    /* Experiment (1):
    */
    // qsort(col, SIZE, sizeof(int), cmp);

    for (int step = 0; step < STEPS; step += 1) {

        /* v grows linearly 0...RAND_MAX in STEPS steps */
        int v = step * (RAND_MAX / (STEPS - 1));

        gettimeofday(&t0, NULL);

        out = 0;
        for (int i = 0; i < SIZE; i += 1) {

```

```
    if (col[i] < v) {
        sv[out] = i;
        out += 1;
    }

    /* Experiment (2):
     * a branch-less copy
     */
    // sv[out] = i;
    // out += (col[i] < v);
}

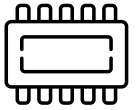
gettimeofday(&t1, NULL);
duration = MICROSECS(t1) - MICROSECS(t0);

selectivity = ((float)out / SIZE) * 100.0;


printf ("%2u (selectivity: %6.2f%%)\t%lu\n",
        step, selectivity, duration);
}

return 0;
}
```


MonetDB: Branch-Less Selection ②



```
❶ for (int i = 0; i < SIZE; i += 1) {  
    if (col[i] < v) {  
        sv[out] = i;  
        out += 1;  
    }  
}
```

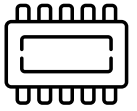


```
❷ for (int i = 0; i < SIZE; i += 1) {  
    sv[out] = i;  
    out += (col[i] < v);  
}  
≡ 1 if predicate satisfied, else 0
```



❷: Only well-predictable loop control flow (**for**) remains.

Mixed-Mode Selection



There is an entire space of possibilities to implement composite predicates (e.g., the conjunction p_1 AND p_2):

- Use branch-less selection via $\text{out} += p_1 \ \& \ p_2$ (note use of C's bit-wise *and* operator $\&$).
- Identify the *more selective*¹ (and thus more predictable) conjunct p_1 , say, then use

```
if ( $p_1$ ) {  
    sv[out] = i;  
    out += ( $p_2$ );  
}
```

¹ **This is important.** Using $\text{if } (p_2) \dots$ instead, where p_2 is unpredictable, immediately ruins the plan.

Demonstrate three alternatives for the implementation of the composite predicate `col < v AND col % 2 = 0`.

See file `mat/mixed-mode-conjunction.c`. Compile via (A no `-O2` here):

```
$ cc -Wall mixed-mode-conjunction.c -o mixed-mode-conjunction
```

- (A) Branch-less selection performs independent of selectivity of both predicates.
- (B) Mixed-mode selection with selective/predictable `col < v` in `if (col < v)` performs better than branch-less as long as `col < v` is very selective (saves work).
- (C) Mixed-mode selection with unpredictable `col % 2 = 0` in `if (col % 2 == 0)` always performs the worst.

Sample run:

```
$ ./mixed-mode-conjunction
  sel      A      mixed B  mixed C
0  0.00%   87ms   28ms    208ms
1  5.00%   87ms   58ms    201ms
2 10.00%   86ms   77ms    198ms
3 15.00%   84ms  101ms    200ms
4 20.00%   86ms  125ms    208ms
5 25.00%   83ms  150ms    205ms    ← selectivity of (col < v) = 50%
6 30.00%   85ms  135ms    195ms
7 35.00%   82ms  122ms    200ms
8 40.00%   82ms  108ms    198ms
9 45.00%   81ms  108ms    209ms
10 50.00%   85ms   87ms    202ms
    ↑
selectivity of col < v AND col % 2 = 0
```

File `mat/mixed-mode-conjunction.c`:

```
/* Demonstrate alternatives for the implementation of
 * conjunctive predicate col < v ^ col % 2 = 0:
 *
 * (A) branch-less selection (via & and +=)
 * (B) mixed mode selection (via if [varying selectivity] and +=)
 * (C) mixed mode selection (via if [unpredictable] and +=)
 */
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <assert.h>

#define MILLISECS(t) ((1000000 * (t).tv_sec + (t).tv_usec) / 1000)
```

```

#define SIZE (16 * 1024 * 1024)
#define STEPS 11

/* comparison of a, b (used in qsort) */
int cmp (const void *a, const void *b)
{
    return *((int*) a) - *((int*) b);
}

int main()
{
    int *col;      /* column vector */
    int *sv;       /* selection vector */

    int out;
    float selectivity;

    struct timeval t0, t1;
    unsigned long duration1, duration2, duration3;

    /* allocate memory */
    col = malloc(SIZE * sizeof(int));
    assert(col);
    sv = malloc(SIZE * sizeof(int));
    assert(sv);

    /* initialize column with (pseudo) random values in interval 0...RAND_MAX */
    srand(42);
    for (int i = 0; i < SIZE; i += 1) {
        col[i] = rand();
    }

    /* Quiz: how will sorting the column affect run time?
    */
    // qsort(col, SIZE, sizeof (*col), cmp);

    printf("\tsel\tA\tmixed B\tmixed C\n");

    for (int step = 0; step < STEPS; step += 1) {

        /* v grows linearly 0...RAND_MAX in STEPS steps */
        int v = step * (RAND_MAX / (STEPS - 1));

        /* ----- alternative A ----- */

        gettimeofday(&t0, NULL);

```

```

    out = 0;
    for (int i = 0; i < SIZE; i += 1) {
        sv[out] = col[i];
        out += ((col[i] < v) & (col[i] % 2 == 0));
    }

    gettimeofday(&t1, NULL);
    duration1 = MILLISECS(t1) - MILLISECS(t0);

    /* ----- alternative B ----- */

    gettimeofday(&t0, NULL);

    out = 0;
    for (int i = 0; i < SIZE; i += 1) {
        if (col[i] < v) {
            sv[out] = col[i];
            out += (col[i] % 2 == 0);
        }
    }

    gettimeofday(&t1, NULL);
    duration2 = MILLISECS(t1) - MILLISECS(t0);

    /* ----- alternative C ----- */

    gettimeofday(&t0, NULL);

    out = 0;
    for (int i = 0; i < SIZE; i += 1) {
        if (col[i] % 2 == 0) {
            sv[out] = col[i];
            out += (col[i] < v);
        }
    }

    gettimeofday(&t1, NULL);
    duration3 = MILLISECS(t1) - MILLISECS(t0);

    selectivity = ((float)out / SIZE) * 100.0;

    printf ("%2u\t5.2f%\t%lums\t%lums\t%lums\n",
            step, selectivity, duration1, duration2, duration3);
}

```

```
} return 0;
```

⚠ File [mat/mixed-mode-selection.c](#) contains even more implementation alternatives, turn this into an assignment?