

# DB 2

---

14 – Query Optimization

Summer 2018

Torsten Grust  
Universität Tübingen, Germany

# 1 | One Query — Millions of Plans

---

**Q:** Given a SQL query  $Q$ , what is *the optimal* (a reasonable)<sup>1</sup> plan to evaluate it? — **A:** It depends:

- Can we **simplify** (flatten, unnest)  $Q$ ?
- How can we **access the tables** referenced in  $Q$ ?
- How do **CPU and (sequential, random) I/O cost** compare?
- What is the **selectivity of the predicates** used in  $Q$ ?
- Which plan **operator implementations** are applicable?
- Can we **regroup/reorder the joins** in  $Q$ ?

<sup>1</sup> Here: focus on reducing the overall query evaluation time. The optimum is, generally, not reached.

## Excerpt of the TPC-H Benchmark (at Scale Factor $SF$ )

---

<u>o_orderkey</u>	<u>o_custkey</u>	<u>o_totalprice</u>	<u>o_clerk</u>	...
$o$	$c$			

orders ( $\approx SF \times 1.5 \times 10^6$  rows)

<u>l_orderkey</u>	<u>l_linenum</u>	<u>l_partkey</u>	<u>l_quantity</u>	<u>l_extendedprice</u>	...
$o$					

lineitem ( $\approx SF \times 6 \times 10^6$  rows)

<u>c_custkey</u>	<u>c_name</u>	<u>c_acctbal</u>	<u>c_nationkey</u>	...
$c$			$n$	

customer ( $\approx SF \times 150000$  rows)

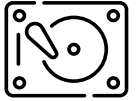
<u>n_nationkey</u>	<u>n_name</u>	<u>n_regionkey</u>	...
$n$		$r$	

nation (25 rows)

<u>r_regionkey</u>	<u>r_name</u>	...
$r$		

region (5 rows)

## Q<sub>14</sub>: Three-Way Join Against a TPC-H Instance



Price and quantity of parts ordered by customer #001:

```
SELECT l.l_partkey, l.l_quantity, l.l_extendedprice
FROM   lineitem AS l JOIN orders AS o      -- { l ⋈ o
      ON (l.l_orderkey = o.o_orderkey)    -- }
      JOIN customer AS c                  -- } ⋈ c
      ON (o.o_custkey = c.c_custkey)      --
WHERE  c.c_name = 'Customer#001';
```

- Above SQL syntax suggests the **join order**  $(l \bowtie o) \bowtie c$ .
- Commutativity and associativity of  $\bowtie$  enable the RDBMS to **reorder** the joins—based on *estimated evaluation costs*.
  - ... unless we insist on the syntactic order. 🙄

Demonstrate the impact of query optimization (force join reordering off) and show a (simple) example of unnesting in the FROM clause.

```
-- 0 Check input tables
-- (this relies on a small TPC-H instance with SF = 0.01, see Week14/live/TPC-H)
```

```
\d lineitem
```

Table "public.lineitem"

Column	Type	Collation	Nullable	Default
l_orderkey	<u>integer</u>		not null	
l_partkey	<u>integer</u>		not null	
l_suppkey	<u>integer</u>		not null	
l_linenum	<u>integer</u>		not null	
l_quantity	<u>numeric(15,2)</u>		not null	
l_extendedprice	<u>numeric(15,2)</u>		not null	
l_discount	<u>numeric(15,2)</u>		not null	
l_tax	<u>numeric(15,2)</u>		not null	
l_returnflag	<u>character(1)</u>		not null	
l_linestatus	<u>character(1)</u>		not null	
l_shipdate	<u>date</u>		not null	
l_commitdate	<u>date</u>		not null	
l_receiptdate	<u>date</u>		not null	
l_shipinstruct	<u>character(25)</u>		not null	
l_shipmode	<u>character(10)</u>		not null	
l_comment	<u>character varying(44)</u>		not null	

**Indexes:**

"lineitem\_pkey" PRIMARY KEY, btree (l\_orderkey, l\_linenum)

**Foreign-key constraints:**

"lineitem\_l\_orderkey\_fkey" FOREIGN KEY (l\_orderkey) REFERENCES orders(o\_orderkey)

"lineitem\_l\_partkey\_fkey" FOREIGN KEY (l\_partkey, l\_suppkey) REFERENCES partsupp(ps\_partkey, ps\_suppkey)

```
\d orders
```

Table "public.orders"

Column	Type	Collation	Nullable	Default
o_orderkey	<u>integer</u>		not null	
o_custkey	<u>integer</u>		not null	
o_orderstatus	<u>character(1)</u>		not null	
o_totalprice	<u>numeric(15,2)</u>		not null	
o_orderdate	<u>date</u>		not null	
o_orderpriority	<u>character(15)</u>		not null	
o_clerk	<u>character(15)</u>		not null	
o_shippriority	<u>integer</u>		not null	

o_comment	character varying(79)		not null	
-----------	-----------------------	--	----------	--

#### Indexes:

"orders\_pkey" PRIMARY KEY, btree (o\_orderkey)

#### Referenced by:

TABLE "lineitem" CONSTRAINT "lineitem\_l\_orderkey\_fkey" FOREIGN KEY (l\_orderkey) REFERENCES orders(o\_orderkey)

\d customer

Table "public.customer"

Column	Type	Collation	Nullable	Default
c_custkey	integer		not null	
c_name	character varying(25)		not null	
c_address	character varying(40)		not null	
c_nationkey	integer		not null	
c_phone	character(15)		not null	
c_acctbal	numeric(15,2)		not null	
c_mktsegment	character(10)		not null	
c_comment	character varying(117)		not null	

#### Indexes:

"customer\_pkey" PRIMARY KEY, btree (c\_custkey)

#### Foreign-key constraints:

"customer\_c\_nationkey\_fkey" FOREIGN KEY (c\_nationkey) REFERENCES nation(n\_nationkey)

-- 2 Evaluate Q<sub>14</sub>, force join reordering/unnesting OFF

```
set join_collapse_limit = 1;
set from_collapse_limit = 1;
```

-- (a) Explicit join order via JOIN ... ON (join\_collapse\_limit)

EXPLAIN (ANALYZE)

```
SELECT l.l_partkey, l.l_quantity, l.l_extendedprice
FROM   lineitem AS l JOIN orders AS o
      ON (l.l_orderkey = o.o_orderkey)
      JOIN customer AS c
      ON (o.o_custkey = c.c_custkey)
WHERE  c.c_name = 'Customer#000000001';
```

QUERY PLAN (c ⋈ (l ⋈ o))
Nested Loop (cost=598.50..3894.86 rows=40 width=17) (actual time=23.063..65.694 rows=35 loops=1) Join Filter: (o.o_custkey = c.c_custkey) Rows Removed by Join Filter: 60140 -> Seq Scan on customer c (cost=0.00..54.75 rows=1 width=4) (actual time=0.022..0.468 rows=1 loops=1)

```

Filter: ((c_name)::text = 'Customer#000000001'::text)
Rows Removed by Filter: 1499
-> Hash Join (cost=598.50..3087.92 rows=60175 width=21) (actual time=13.785..56.956 rows=60175 loops=1)
    Hash Cond: (l.l_orderkey = o.o_orderkey)
    -> Seq Scan on lineitem l (cost=0.00..1729.75 rows=60175 width=21) (actual time=0.011..10.405 rows=60175 loops=1)
    -> Hash (cost=411.00..411.00 rows=15000 width=8) (actual time=13.746..13.746 rows=15000 loops=1)
        Buckets: 16384 Batches: 1 Memory Usage: 714kB
    -> Seq Scan on orders o (cost=0.00..411.00 rows=15000 width=8) (actual time=0.013..6.647 rows=15000 loops=1)
Planning time: 1.115 ms
Execution time: 65.792 ms ◀ slow

```

-- (b) Prescribed join order via subquery nesting in the FROM clause (from\_collapse\_limit)

EXPLAIN (ANALYZE)

```

SELECT lo.l_partkey, lo.l_quantity, lo.l_extendedprice
FROM (SELECT l.l_partkey, l.l_quantity, l.l_extendedprice, o.o_custkey
      FROM lineitem AS l, orders AS o
      WHERE l.l_orderkey = o.o_orderkey) AS lo,
      customer AS c
WHERE c.c_name = 'Customer#000000001'
AND lo.o_custkey = c.c_custkey;

```

#### QUERY PLAN (c ⋈ (l ⋈ o))

```

Nested Loop (cost=598.50..3894.86 rows=40 width=17) (actual time=18.125..54.639 rows=35 loops=1)
  Join Filter: (o.o_custkey = c.c_custkey)
  Rows Removed by Join Filter: 60140
  -> Seq Scan on customer c (cost=0.00..54.75 rows=1 width=4) (actual time=0.021..0.319 rows=1 loops=1)
      Filter: ((c_name)::text = 'Customer#000000001'::text)
      Rows Removed by Filter: 1499
  -> Hash Join (cost=598.50..3087.92 rows=60175 width=21) (actual time=10.879..48.201 rows=60175 loops=1)
      Hash Cond: (l.l_orderkey = o.o_orderkey)
      -> Seq Scan on lineitem l (cost=0.00..1729.75 rows=60175 width=21) (actual time=0.011..8.804 rows=60175 loops=1)
      -> Hash (cost=411.00..411.00 rows=15000 width=8) (actual time=10.842..10.842 rows=15000 loops=1)
          Buckets: 16384 Batches: 1 Memory Usage: 714kB
      -> Seq Scan on orders o (cost=0.00..411.00 rows=15000 width=8) (actual time=0.011..5.293 rows=15000 loops=1)
Planning time: 0.926 ms
Execution time: 54.733 ms ◀ slow

```

-- Plan shape and cardinalities:

```

--
--      | 35
--      ⋈
--  1 /   \ 60175
--    σ     ⋈

```

```
--      |  |  \
--    1500 c  1  o  15000
--      |  |  \
--      60175
```

```
-- 2 Re-evaluate Q14, with join reordering/unnesting enabled
```

```
reset join_collapse_limit;
reset from_collapse_limit;
```

```
EXPLAIN (ANALYZE)
```

```
SELECT l1.partkey, l1.quantity, l1.extendedprice
FROM   lineitem AS l1 JOIN orders AS o
      ON (l1.orderkey = o.orderkey)
JOIN   customer AS c
      ON (o.custkey = c.custkey)
WHERE  c.c_name = 'Customer#000000001';
```

QUERY PLAN ((c ⋈ o) ⋈<sub>idx 1</sub> l)

```
Nested Loop (cost=0.29..660.70 rows=40 width=17) (actual time=1.660..7.058 rows=35 loops=1)
-> Nested Loop (cost=0.00..653.25 rows=10 width=4) (actual time=1.633..6.905 rows=9 loops=1)
    Join Filter: (o.o_custkey = c.c_custkey)
    Rows Removed by Join Filter: 14991
    -> Seq Scan on customer c (cost=0.00..54.75 rows=1 width=4) (actual time=0.023..0.367 rows=1 loops=1)
        Filter: ((c_name)::text = 'Customer#000000001'::text)
        Rows Removed by Filter: 1499
    -> Seq Scan on orders o (cost=0.00..411.00 rows=15000 width=8) (actual time=0.011..3.429 rows=15000 loops=1)
-> Index Scan using lineitem_pkey on lineitem l1 (cost=0.29..0.71 rows=4 width=21) (actual time=0.009..0.012 rows=4 loops=9)
    Index Cond: (l1_orderkey = o.o_orderkey)
Planning time: 1.185 ms
Execution time: 7.145 ms ◀ fast
```

```
-- Plan shape and cardinalities:
```

```
--      | 35
--      | ⋈
--      | / \
--      9 /   \
--      | ⋈     |
--      | / \   |
--      1 /   \ |
--      | σ     |
--    1500 c  o  1  60175
--      |  |  \
--      15000
```



```

EXPLAIN (ANALYZE)
SELECT lo.l_partkey, lo.l_quantity, lo.l_extendedprice
FROM (SELECT l.l_partkey, l.l_quantity, l.l_extendedprice, o.o_custkey
      FROM lineitem AS l, orders AS o
      WHERE l.l_orderkey = o.o_orderkey) AS lo,
      customer AS c
WHERE c.c_name = 'Customer#000000001'
AND lo.o_custkey = c.c_custkey;

```


QUERY PLAN ((c ⋈ o) ⋈<sup>idx</sup> l)

```

Nested Loop (cost=0.29..660.70 rows=40 width=17) (actual time=1.616..6.295 rows=35 loops=1)
-> Nested Loop (cost=0.00..653.25 rows=10 width=4) (actual time=1.582..6.177 rows=9 loops=1)
    Join Filter: (o.o_custkey = c.c_custkey)
    Rows Removed by Join Filter: 14991
    -> Seq Scan on customer c (cost=0.00..54.75 rows=1 width=4) (actual time=0.023..0.324 rows=1 loops=1)
        Filter: ((c_name)::text = 'Customer#000000001'::text)
        Rows Removed by Filter: 1499
    -> Seq Scan on orders o (cost=0.00..411.00 rows=15000 width=8) (actual time=0.011..2.891 rows=15000 loops=1)
-> Index Scan using lineitem_pkey on lineitem l (cost=0.29..0.71 rows=4 width=21) (actual time=0.009..0.010 rows=4 loops=9)
    Index Cond: (l_orderkey = o.o_orderkey)
Planning time: 1.214 ms
Execution time: 6.389 ms ◀ fast

```

## 2 | Pre-Processing: Query Normalization



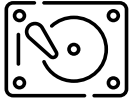
Transform the input SQL query such that it features  
**SELECT-FROM-WHERE** (SFW) blocks of the following shape:

```
SELECT [DISTINCT] e, ..., e
FROM    Δ, ..., Δ
[ WHERE  p AND ... AND p ]
[ GROUP BY g, ..., g
  [ HAVING p AND ... AND p ] ]
[ ORDER BY o, ..., o ]
[ OFFSET n ]
[ LIMIT  m ]
```

-- Δ ≡ base table or (*query*)  
-- p ≡ predicate in DNF  
-- { e, p, g, o ≡  
-- atomic expression or  
-- scalar (*subquery*)  
-- { n, m ≡ integer literal  
-- }

- Query clauses in [...] may be missing.


### 3 | Pre-Processing: Query Unnesting



**Nested SQL queries** suggest a (naïve, inefficient) nested-loop-style evaluation strategy. Consider:

<pre>SELECT c.c_name FROM   customer AS c,       ⚠ { (SELECT n.n_nationkey, n.n_name           \ FROM   nation AS n) AS t WHERE  c.c_nationkey = t.n_nationkey       AND strpos(c.c_address, t.n_name) &gt; 0</pre>	<pre>SELECT o.o_orderkey FROM   orders AS o WHERE  o.o_custkey IN       ⚠ { (SELECT c.c_custkey           \ FROM   customer AS c             WHERE  c.c_name = '...')</pre>
---	---

- 💡 If possible, **unnest** ⚠ queries and “inline” into parent query  $\Rightarrow$  ⚠ can participate in join reordering.



Perform **query unnesting** on the level of

- the operator-based plan representation of the query,<sup>2</sup> or
- the internal AST representation of SQL. Re ②:

<b>SELECT</b> $e_1$ <b>FROM</b> $q_1, \dots, q_i$ <b>WHERE</b> $p_1$ <b>AND</b> $e_2$ <b>IN</b> ( <b>SELECT</b> $e_3$ <b>FROM</b> $q_{i+1}, \dots, q_n$ <b>WHERE</b> $p_3$ )	$\cong$	<b>SELECT DISTINCT</b> $e_1$ <b>FROM</b> $q_1, \dots, q_i, q_{i+1}, \dots, q_n$ <b>WHERE</b> $p_1$ <b>AND</b> $e_2 = e_3$ <b>AND</b> $p_3$
---	---------	--

\* Precondition:  $e_1$  is key in the left-hand side query

<sup>2</sup> See *Unnesting Arbitrary Queries*, Thomas Neumann, Alfons Kemper. BTW 2015, Hamburg, Germany.

Demonstrate subquery unnesting for nesting in the ❶ **FROM** clause and nesting in the ❷ **WHERE** clause.

- ❶ Unnesting

- N.B.: Rename **t** → **n** as required:

```
SELECT c.c_name
FROM   customer AS c, nation AS n
WHERE  c.c_nationkey = n.n_nationkey
AND    strpos(c.c_address, n.n_name) > 0;
```

- ❷ Unnesting (**IN** → join + **DISTINCT**)

- ⚠ **DISTINCT** actually *not* needed in this special case since for each **o** there will be at most one join partner in **customer** (**c\_custkey** is key).

```
SELECT DISTINCT o.o_orderkey
FROM   orders AS o, customer AS c
WHERE  c.c_name = 'Customer#0000000001'
AND    o.o_custkey = c.c_custkey;
```

-- ❶ Nesting in the FROM clause: inlining

```
EXPLAIN (COSTS false)
SELECT c.c_name
FROM   customer AS c,
      (SELECT n.n_nationkey, n.n_name
       FROM   nation AS n) AS t
WHERE  c.c_nationkey = t.n_nationkey
AND    strpos(c.c_address, t.n_name) > 0;
```

QUERY PLAN
Hash Join Hash Cond: (c.c_nationkey = n.n_nationkey) Join Filter: (strpos((c.c_address)::text, (n.n_name)::text) > 0) -> Seq Scan on customer c -> Hash -> Seq Scan on nation n

-- Manual unnesting: identical plan

```
EXPLAIN (COSTS false)
```

```

SELECT c.c_name
FROM   customer AS c, nation AS n
WHERE  c.c_nationkey = n.n_nationkey
AND    strpos(c.c_address, n.n_name) > 0;

```

#### QUERY PLAN

```

Hash Join
Hash Cond: (c.c_nationkey = n.n_nationkey)
Join Filter: (strpos((c.c_address)::text, (n.n_name)::text) > 0)
-> Seq Scan on customer c
-> Hash
    -> Seq Scan on nation n

```

-- 2 Nesting in the WHERE clause: IN ⇒ semijoin

```

EXPLAIN (COSTS false)
SELECT o.o_orderkey
FROM   orders AS o
WHERE  o.o_custkey IN
      (SELECT c.c_custkey
       FROM   customer AS c
       WHERE  c.c_name = 'Customer#000000001');

```

#### QUERY PLAN

```

Nested Loop
Join Filter: (o.o_custkey = c.c_custkey) ← will find at most one match for o ⇒ no Unique required
-> Seq Scan on customer c
    Filter: ((c_name)::text = 'Customer#000000001'::text)
-> Seq Scan on orders o

```

-- Variant of 2: do not compare with key 'c\_custkey' ⇒ optimizer uses hash-based duplicate elimination to ensure that each o will find at most one join partner

```

EXPLAIN (COSTS false)
SELECT o.o_orderkey
FROM   orders AS o
WHERE  o.o_clerk IN
      (SELECT c.c_name
       FROM   customer AS c);

```

## QUERY PLAN

### Hash Join

Hash Cond: (o.o\_clerk = (c.c\_name)::bpchar)

-> Seq Scan on orders o

-> Hash

-> HashAggregate

Group Key: (c.c\_name)::bpchar

-> Seq Scan on customer c

will find at most one match for o → no Unique required

} build table of  
unique customer  
names

-- Variant of 2: switch off hash join → optimizer uses (Merge) Semi Join to ensure that each o is matched with at most one join partner

set enable\_hashjoin = off;

EXPLAIN (COSTS false)

SELECT o.o\_orderkey

FROM orders AS o

WHERE o.o\_clerk IN

(SELECT c.c\_name

FROM customer AS c);

## QUERY PLAN

### Merge Semi Join

Merge Cond: (o.o\_clerk = (c.c\_name)::bpchar)

-> Sort

Sort Key: o.o\_clerk

-> Seq Scan on orders o


-> Sort

Sort Key: c.c\_name USING <

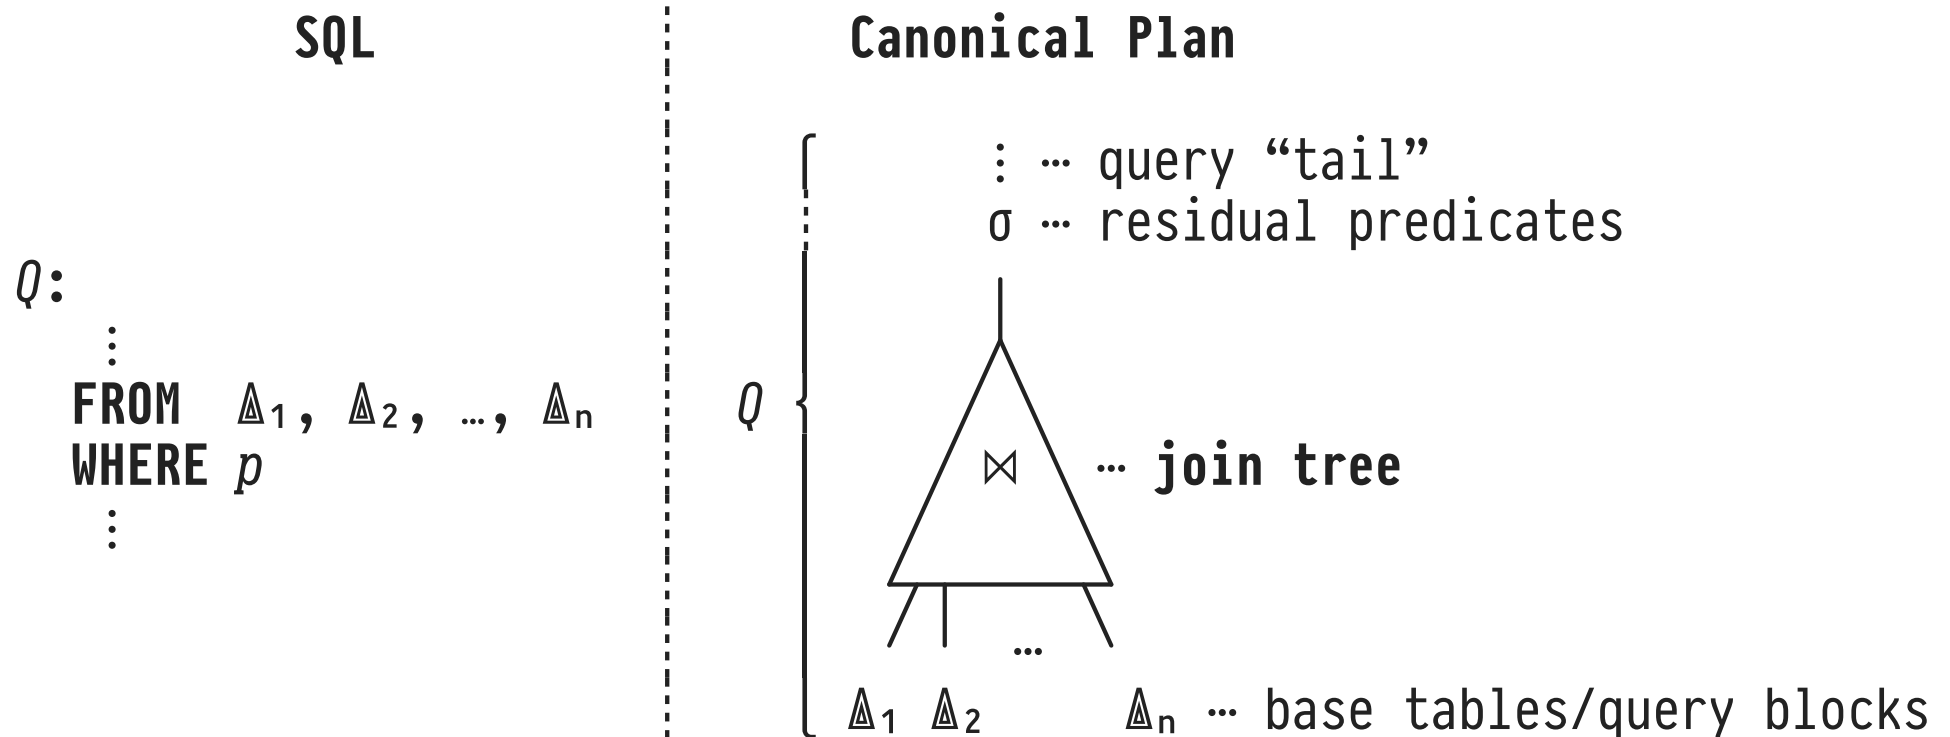
-> Seq Scan on customer c

set enable\_hashjoin = on;

## 4 | Join Tree Optimization




Processing a SQL query  $Q$  starts out with its **FROM** and **WHERE** clauses which describe a **join tree** over  $Q$ 's inputs:






# Join Tree Optimization




Given  $n$  join inputs, the number of possible **join tree shapes** is *huge*. Consider  $n = 3$ :




- Shapes based on associativity and commutativity of  $\Join$ .

# How Many Possible Join Trees are There?



1. A join of  $n+1$  inputs  $\Delta$  requires  $n$  binary joins. The root  $\bowtie$  combines subtrees of  $k$  and  $n-k-1$  joins ( $0 \leq k \leq n-1$ ):<sup>3</sup>



# of join tree shapes:

$$C_n = \sum_{k=0}^{n-1} C_k \times C_{n-k-1}$$

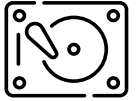
2. Orderings of the  $\Delta$  at the join tree leaf level:  $(n+1)!$ .
3. Join algorithm choices ( $a$  available algorithms):  $a^n$ .

<sup>3</sup>  $C_n$  are the *Catalan numbers*, the number of ordered binary trees with  $n+1$  leaves.  $C_0 = 1$ .

Catalan numbers  $\equiv$  possible tree shapes:

- $C_0 = 1$  (single leaf node, no inner nodes)
- $C_1 = 1$  (single inner node)
- $C_2 = C_0 \times C_1 + C_1 \times C_0 = 1 + 1 = 2$

## How Many Possible Join Trees are There?



Number of possible join trees given  $n$  binary joins with  $a = 3$  implementation choices:

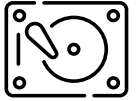
# of $\triangle (n+1)$	$C_n$	# of join trees
2	1	6
3	2	108
4	5	3240
5	14	136080
6	42	7384320
7	132	484989120
8	429	37829151360
9	1430	3404623622400
10	4862	347271609484800

- A search space of this size is impossible to fully explore for any query optimizer.

- 347271609484800  $\equiv$  347 trillion 271 billion 609 million 484 thousand 800


# Join Plan Generation Through Dynamic Programming

---



- **Problem:** Find optimal query plan  $opt[\{\Delta_1, \dots, \Delta_n\}]$  that joins  $n$  inputs  $\Delta_1, \dots, \Delta_n$ .
  1. **Iteration 1:** For each  $\Delta_j$ , find and memorize **best 1-input plan**  $opt[\{\Delta_j\}]$  that accesses  $\Delta_j$  only.
  2. **Iteration  $k > 1$ :** Find and memorize **best  $k$ -input plans** that join  $k \leq n$  inputs by combining (for  $1 \leq i < k$ )
    - the best  $i$ -input plans and  $\setminus$  simple lookups in
    - the best  $(k-i)$ -input plans.  $\int opt[\cdot]$  memo 👍

# Bottom-Up Dynamic Programming ( $n = 3$ )



**$k$  Possible  $k$ -input Access/Join Plans** if  $\Delta_i$  is complex


**1**  $opt[\{\Delta_1\}] \leftarrow prune(\{Seq\ Scan\ \Delta_1, Index\ Scan\ \Delta_1, Bitmap\ Scan\ \Delta_1, \overbrace{\Delta_1}^{complex}\})$   
 $opt[\{\Delta_2\}] \leftarrow prune(\{Seq\ Scan\ \Delta_2, Index\ Scan\ \Delta_2, Bitmap\ Scan\ \Delta_2, \Delta_2\})$   
 $opt[\{\Delta_3\}] \leftarrow prune(\{Seq\ Scan\ \Delta_3, Index\ Scan\ \Delta_3, Bitmap\ Scan\ \Delta_3, \Delta_3\})$

**2**  $opt[\{\Delta_1, \Delta_2\}] \leftarrow prune(opt[\{\Delta_1\}] \otimes opt[\{\Delta_2\}])$   
 $opt[\{\Delta_1, \Delta_3\}] \leftarrow prune(opt[\{\Delta_1\}] \otimes opt[\{\Delta_3\}])$   
 $opt[\{\Delta_2, \Delta_3\}] \leftarrow prune(opt[\{\Delta_2\}] \otimes opt[\{\Delta_3\}])$

**3**  $opt[\{\Delta_1, \Delta_2, \Delta_3\}] \leftarrow prune(opt[\{\Delta_1\}] \otimes opt[\{\Delta_2, \Delta_3\}] \cup$   
 $\quad opt[\{\Delta_2\}] \otimes opt[\{\Delta_1, \Delta_3\}] \cup$   
 $\quad opt[\{\Delta_3\}] \otimes opt[\{\Delta_1, \Delta_2\}] )$

$prune(P) \equiv$  best (= minimal cost + interestingly ordered) plans in set  $P$


$l \otimes r \equiv \{l \bowtie^{n1} r, r \bowtie^{n1} l, l \bowtie^{mj} r, r \bowtie^{mj} l, l \bowtie^{hj} r, r \bowtie^{hj} l\}$



- **Access plan choices** (*access(·)*):
  - Consider sequential/index scans if  $\Delta$  is a base table, otherwise simply consume  $\Delta$ 's rows.
- **Join plan choices** (*\_  $\Join$  \_*):
  - Considers all viable join algorithms (given  $\theta$ , available indexes, ...) and left/right input orders.
- **Principle of Optimality** (*prune(·)*): A globally optimal plan is built from optimal subplans. Thus:
  - 💡 For each subset of  $\{\Delta_1, \dots, \Delta_n\}$ , memorize in *opt[·]*
    1. ... its overall best plan and
    2. ... its best plan satisfying each **interesting order**.



## (Bushy) Join Plan Generation: Pseudo Code



```
JoinPlan( $\{\Delta_1, \dots, \Delta_n\}$ ):  
  foreach  $p \in \{\Delta_1, \dots, \Delta_n\}$  } 1-input plans  
  [  $opt[\{p\}] \leftarrow prune(access(p))$ ;  
  
  for  $k$  in  $2, \dots, n$  }  $k$ -input plans  
  [ foreach  $S \subseteq \{\Delta_1, \dots, \Delta_n\}$  with  $|S| = k$  } enumerate subsets  
  [    $opt[S] \leftarrow \phi$ ;  
  [   foreach  $T \subset S$  with  $T \neq \phi$   $\Join^a \setminus$   
  [     [  $opt[S] \leftarrow opt[S] \cup \{ opt[T] \Join^a opt[S \setminus T] \}$ ;  
  [     [  $opt[S] \leftarrow prune(opt[S])$ ;  
  
  return  $opt[\{\Delta_1, \dots, \Delta_n\}]$ ;
```

- $access(\cdot)$ ,  $prune(\cdot)$  defined as above,  
  $\Join^a \setminus$  builds all join algorithm choices ( $a \in \{nl, mj, hj\}$ ).


- Enumerate all non-empty true subsets of  $S$  (bit set representation in C):

```
T = S & ~S;
do {
    /* perform operation on T */
    T = S & (T ^ S);
} while (T != S);
```




Building the  $opt[\cdot]$  and pruning plans early leads to a compact representation of the explored search space. Example for  $n = 4$ , all  $\mathbb{A}_i$  base tables (i.e.,  $access(\mathbb{A}_i)$  considers Seq Scan, Index Scan, Bitmap Scan):

$opt[\{\mathbb{A}_1\}] = \text{prune}(\text{access}(\mathbb{A}_1)) \leftarrow 3 \text{ plans considered}$ $opt[\{\mathbb{A}_2\}] = \dots$ $opt[\{\mathbb{A}_3\}] = \dots$ $opt[\{\mathbb{A}_4\}] = \dots$	$\left. \vphantom{\begin{array}{l} opt[\{\mathbb{A}_1\}] = \text{prune}(\text{access}(\mathbb{A}_1)) \leftarrow 3 \text{ plans considered} \\ opt[\{\mathbb{A}_2\}] = \dots \\ opt[\{\mathbb{A}_3\}] = \dots \\ opt[\{\mathbb{A}_4\}] = \dots \end{array}} \right\} 4 \times 3 = 12 \text{ plans considered,}$ only 4 plans memorized
$opt[\{\mathbb{A}_1, \mathbb{A}_2\}] = \text{prune}(opt[\{\mathbb{A}_1\}] \otimes opt[\{\mathbb{A}_2\}]) \leftarrow 6 \text{ plans considered}$ $opt[\{\mathbb{A}_1, \mathbb{A}_3\}] = \dots$ $opt[\{\mathbb{A}_1, \mathbb{A}_4\}] = \dots$ $opt[\{\mathbb{A}_2, \mathbb{A}_3\}] = \dots$ $opt[\{\mathbb{A}_2, \mathbb{A}_4\}] = \dots$ $opt[\{\mathbb{A}_3, \mathbb{A}_4\}] = \dots$	$\left. \vphantom{\begin{array}{l} opt[\{\mathbb{A}_1, \mathbb{A}_2\}] = \text{prune}(opt[\{\mathbb{A}_1\}] \otimes opt[\{\mathbb{A}_2\}]) \leftarrow 6 \text{ plans considered} \\ opt[\{\mathbb{A}_1, \mathbb{A}_3\}] = \dots \\ opt[\{\mathbb{A}_1, \mathbb{A}_4\}] = \dots \\ opt[\{\mathbb{A}_2, \mathbb{A}_3\}] = \dots \\ opt[\{\mathbb{A}_2, \mathbb{A}_4\}] = \dots \\ opt[\{\mathbb{A}_3, \mathbb{A}_4\}] = \dots \end{array}} \right\} 6 \times 6 = 36 \text{ plans considered,}$ only 6 plans memorized
$opt[\{\mathbb{A}_1, \mathbb{A}_2, \mathbb{A}_3\}] = \left. \begin{array}{l} opt[\{\mathbb{A}_1\}] \otimes opt[\{\mathbb{A}_2, \mathbb{A}_3\}] \cup \\ opt[\{\mathbb{A}_2\}] \otimes opt[\{\mathbb{A}_1, \mathbb{A}_3\}] \cup \\ opt[\{\mathbb{A}_3\}] \otimes opt[\{\mathbb{A}_1, \mathbb{A}_2\}] \end{array} \right\} 18 \text{ plans considered}$ $opt[\{\mathbb{A}_1, \mathbb{A}_3, \mathbb{A}_4\}] = \dots$ $opt[\{\mathbb{A}_1, \mathbb{A}_2, \mathbb{A}_4\}] = \dots$ $opt[\{\mathbb{A}_2, \mathbb{A}_3, \mathbb{A}_4\}] = \dots$	$\left. \vphantom{\begin{array}{l} opt[\{\mathbb{A}_1, \mathbb{A}_2, \mathbb{A}_3\}] = \left. \begin{array}{l} opt[\{\mathbb{A}_1\}] \otimes opt[\{\mathbb{A}_2, \mathbb{A}_3\}] \cup \\ opt[\{\mathbb{A}_2\}] \otimes opt[\{\mathbb{A}_1, \mathbb{A}_3\}] \cup \\ opt[\{\mathbb{A}_3\}] \otimes opt[\{\mathbb{A}_1, \mathbb{A}_2\}] \end{array} \right\} 18 \text{ plans considered} \\ opt[\{\mathbb{A}_1, \mathbb{A}_3, \mathbb{A}_4\}] = \dots \\ opt[\{\mathbb{A}_1, \mathbb{A}_2, \mathbb{A}_4\}] = \dots \\ opt[\{\mathbb{A}_2, \mathbb{A}_3, \mathbb{A}_4\}] = \dots \end{array}} \right\} 4 \times 18 = 72 \text{ plans considered,}$ only 4 plans memorized
$opt[\{\mathbb{A}_1, \mathbb{A}_2, \mathbb{A}_3, \mathbb{A}_4\}] = \left. \begin{array}{l} opt[\{\mathbb{A}_1\}] \otimes opt[\{\mathbb{A}_2, \mathbb{A}_3, \mathbb{A}_4\}] \cup \\ opt[\{\mathbb{A}_2\}] \otimes opt[\{\mathbb{A}_1, \mathbb{A}_3, \mathbb{A}_4\}] \cup \\ opt[\{\mathbb{A}_3\}] \otimes opt[\{\mathbb{A}_1, \mathbb{A}_2, \mathbb{A}_4\}] \cup \\ opt[\{\mathbb{A}_4\}] \otimes opt[\{\mathbb{A}_1, \mathbb{A}_2, \mathbb{A}_3\}] \cup \\ opt[\{\mathbb{A}_1, \mathbb{A}_2\}] \otimes opt[\{\mathbb{A}_3, \mathbb{A}_4\}] \cup \\ opt[\{\mathbb{A}_1, \mathbb{A}_3\}] \otimes opt[\{\mathbb{A}_2, \mathbb{A}_4\}] \cup \\ opt[\{\mathbb{A}_1, \mathbb{A}_4\}] \otimes opt[\{\mathbb{A}_2, \mathbb{A}_3\}] \end{array} \right\}$	42 plans considered, only 1 plan memorized
$\Sigma$ 162 plans considered 15 plans memorized	

## Reducing the Search Space




- Avoid generating costly **Cartesian products**: don't form joins between inputs w/o join predicate (`_  $\theta$  _ = true`).
- Generate **left-deep** join plans only: right join input (NL $\bowtie$ : inner input) is a scan over base table  $T$ .
  - Admits use of Index Nested Loop Join.
  - Straightforward Volcano-style execution (reset inner).



## 5 | Estimating Plan Cost

---



The query optimizer explores the vast plan search space to find the **optimal** (“best”, “cheapest”) plan.

- Typically, RDBMSs measure **plan cost** in terms of *total execution time* (time until last result row delivered).
- These total plan costs are **estimated before** plan execution begins (`EXPLAIN: ... cost=c1..c2← ...`).
- A **cost model**—measured in abstract “space\$”—reflects the true costs (measured in *ms*, CPU time, # I/O ops, ...) of plans  $p_1$ ,  $p_2$ :

$$\text{space}\$(p_1) < \text{space}\$(p_2) \Rightarrow \text{true cost}(p_1) < \text{true cost}(p_2)$$

# PostgreSQL: Plan Cost



**EXPLAIN** shows estimated costs (unit: space\$) and cardinalities (# of rows):

QUERY PLAN		
	startup cost	total cost
Hash Join		
⋮		
	(cost=299.00..15443.31	rows=505183 width=50)
		cardinality

- **run cost**  $\stackrel{\text{def}}{=} \text{total cost} - \text{startup cost}$ <sup>4</sup> (not shown).
- Optimizer decisions are based on estimated **total cost**.

<sup>4</sup> To implement `set enable_<op> = off`, PostgreSQL sets the operator's **startup cost** to  $10^9$  ( $\equiv \infty$ ).

Demonstrate that the PostgreSQL space\$-based cost model does not try to estimate the true cost of plan evaluation:

-- 1 Check input table

\d ternary

Table "public.ternary"

Column	Type	Collation	Nullable	Default
a	<u>integer</u>		<b>not</b> null	
b	<u>text</u>		<b>not</b> null	
c	<u>double</u> precision			

-- 2 Plan/evaluate two queries with same space\$ cost but \*wildly\*  
-- different true cost

**EXPLAIN** (VERBOSE, ANALYZE)  
**SELECT** t.a::bigint + 1 -- same data type as used by \_!  
**FROM** ternary AS t;


QUERY PLAN
Seq Scan on public.ternary t (cost=0.00..25.00 rows=1000 width=8) (actual time=0.060..0.554 rows=1000 loops=1) Output: ((a)::bigint + 1) Planning time: 0.063 ms Execution time: 0.756 ms ◀ fast

**EXPLAIN** (VERBOSE, ANALYZE)  
**SELECT** t.a!  
**FROM** ternary AS t;

QUERY PLAN
Seq Scan on public.ternary t (cost=0.00..25.00 rows=1000 width=32) (actual time=0.016..1535.795 rows=1000 loops=1) Output: ((a)::bigint !) Planning time: 0.039 ms Execution time: 1536.038 ms ◀ slow

# Cost Model Configuration

---




Model Configuration	Default	Description
<code>seq_page_cost</code>	1.0	I/O cost of one sequential page access
<code>random_page_cost</code>	4.0	I/O cost of one random page access
<code>cpu_tuple_cost</code>	0.01	CPU cost to process a heap file row
<code>cpu_index_tuple_cost</code>	0.005	CPU cost to process an index leaf entry
<code>cpu_operator_cost</code>	0.0025	CPU function/operator evaluation cost
<code>parallel_tuple_cost</code>	0.1	Cost of passing one row worker→leader
<code>parallel_setup_cost</code>	1000.0	Shared memory setup cost

- Parameters are configurable:
  - Seek cost, thus `random_page_cost` » `seq_page_cost`. But...
  - ... if DB fits in RAM, `random_page_cost` = `seq_page_cost` may be more appropriate.



## Cost of Seq Scan ①




Given an occurrence of **Seq Scan** with arguments

- *in*: input table,
- *pred*: (optional) filter predicate on *in*,
- *expr*: **SELECT** clause expression(s),

how does PostgreSQL derive *startup\_cost* and *total\_cost*?

<i>in</i>	QUERY PLAN	<i>total_cost</i>
Seq Scan on public.indexed i	(cost=0.00..22.75 rows=100 width=4)	
Output: (a + 1)	◀ <i>expr</i>	▲
Filter: (i.a <= 100)	◀ <i>pred</i>	▲
	<i>startup_cost</i>	<i>#rows(out)</i>

## Cost of Seq Scan ②



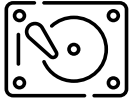
Cost calculation depends on the following parameters, mostly available in PostgreSQL's internal `pg_*` meta data tables:

Parameter	Description	Available as...
<code>#rows(<i>in</i>)</code>	# rows (cardinality) of table <i>in</i>	<code>pg_class.reltuples</code>
<code>#pages(<i>in</i>)</code>	# pages in heap file of <i>in</i>	<code>pg_class.relpages</code>
<code>sel(<i>pred</i>)</code>	selectivity of filter <i>pred</i> <sup>5</sup>	see below

- Meta data like `#rows(in)`, `#pages(in)` and others are updated whenever the system performs an `ANALYZE` run on table *in*.
- Predicate selectivity `sel(pred)` is estimated based on sampled table data and the syntactic structure of *pred*.

<sup>5</sup> `sel(pred)`  $\in \{0, \dots, 1\}$  with `sel(pred) = 0`  $\equiv$  no row satisfies filter *pred*.

## Cost of Seq Scan ③



$\nwarrow$  typically = 0  $\searrow$

**startup\_cost**  $\stackrel{\text{def}}{=} \text{startup\_cost}(\text{pred}) + \text{startup\_cost}(\text{expr})$

**cpu\_run\_cost**  $\stackrel{\text{def}}{=} \underbrace{\#rows(in) \times (\text{cpu\_tuple\_cost} + \text{run\_cost}(\text{pred}))}_{\text{decode heap row}} + \underbrace{\#rows(in) \times \text{sel}(\text{pred}) \times \text{run\_cost}(\text{expr})}_{\text{evaluate filter}}$

$\quad \quad \quad = \#rows(out) \quad \quad \quad \text{evaluate SELECT clause}$

**disk\_run\_cost**  $\stackrel{\text{def}}{=} \underbrace{\#pages(in) \times \text{seq\_page\_cost}}_{\text{sequentially read entire input heap file}}$

**total\_cost**  $\stackrel{\text{def}}{=} \text{startup\_cost} + \underbrace{\text{cpu\_run\_cost} + \text{disk\_run\_cost}}_{= \text{run\_cost}}$

Demonstrate the cost derivation for Seq Scan on table `indexed`.

-- ❶ Set up input table

```
DROP TABLE IF EXISTS indexed;
CREATE TABLE indexed (a int PRIMARY KEY, b text, c numeric(3,2));
```

```
INSERT INTO indexed(a,b,c)
  SELECT i, md5(i::text), sin(i)
  FROM   generate_series(1,1000000) AS i;
```

```
ALTER INDEX indexed_pkey RENAME TO indexed_a;
ANALYZE indexed;
```

\d indexed

Table "public.indexed"

Column	Type	Collation	Nullable	Default
a	integer		not null	
b	text			
c	numeric(3,2)			

Indexes:

"indexed\_a" PRIMARY KEY, btree (a)

-- ❷ Obtain meta data about table indexed

```
SELECT reltuples AS "#rows(indexed)", relpages AS "#pages(indexed)"
FROM   pg_class
WHERE  relname = 'indexed';
```

#rows(indexed)	#pages(indexed)
1e+06	9346

-- ❸ Simple Seq Scan (no filter)

```
EXPLAIN VERBOSE
SELECT i.a
FROM   indexed AS i;
```

QUERY PLAN

Seq Scan on public.indexed i (cost=0.00..19346.00 rows=1000000 width=4) Output: a
--

```
- startup_cost = startup_cost(pred) + startup_cost(expr) = 0 + 0
               = 0.00 ✓

- cpu_run_cost = #rows(indexed) × (cpu_tuple_cost + run_cost(pred))
               + #rows(indexed) × sel(pred) × run_cost(expr)
               = 106 × (0.01 + 0)
               + 106 × 1.0 × 0
               = 10000.00

- disk_run_cost = #pages(indexed) × seq_page_cost
               = 9346 × 1.0
               = 9346.00

- total_cost   = startup_cost + cpu_run_cost + disk_run_cost
               = 0.00 + 10000.00 + 9346.00
               = 19346.00 ✓
```

-- 4 Simple Seq Scan (no filter but "complex" SELECT clause expr)

```
EXPLAIN VERBOSE
SELECT i.a * 2 + 1
FROM indexed AS i;
```

QUERY PLAN
------------

Seq Scan on public.indexed i (cost=0.00..24346.00 rows=1000000 width=4) Output: ((a * 2) + 1)
--

```
- startup_cost = startup_cost(pred) + startup_cost(expr) = 0 + 0
               = 0.00 ✓

- cpu_run_cost = #rows(indexed) × (cpu_tuple_cost + run_cost(pred))
               + #rows(indexed) × sel(pred) × run_cost(expr)
               = 106 × (0.01 + 0)
               + 106 × 1.0 × 2 × 0.0025
               = 15000.00
               ▲
               run_cost(expr) = 2 × cpu_operator_cost: . * ., . + .

- disk_run_cost = #pages(indexed) × seq_page_cost
               = 9346 × 1.0
```

= 9346.00

```
- total_cost = startup_cost + cpu_run_cost + disk_run_cost
              = 0.00 + 15000.00 + 9346.00
              = 24346.00 ✓
```

-- 5 Simple Seq Scan (with filter and SELECT clause expression)

```
-- enforce Seq Scan
set index_scan = off;
set bitmap_scan = off;
```

```
EXPLAIN VERBOSE
SELECT i.a * 2 + 1
FROM indexed AS i
WHERE i.a <= 100000;
```

QUERY PLAN
Seq Scan on public.indexed i (cost=0.00..22342.17 rows=99235 width=4) Output: ((a * 2) + 1) Filter: (i.a <= 100000)

- Estimated #rows(out) = 99325  $\Rightarrow$   $\text{sel}(i.a \leq 100000) = 99325/10^6 = 0.099325$

```
- startup_cost = startup_cost(pred) + startup_cost(expr) = 0 + 0
                = 0.00 ✓
```

```
- cpu_run_cost = #rows(indexed) × (cpu_tuple_cost + run_cost(pred))
                + #rows(indexed) × sel(pred) × run_cost(expr)
                = 10^6 × (0.01 + 0.0025)      ─ run_cost(pred) = cpu_operator_cost: . <= .
                + 10^6 × 0.099325 × 2 × 0.0025 ─ run_cost(expr) = 2 × cpu_operator_cost: . * ., . + .
                = 12996.625
```

```
- disk_run_cost = #pages(indexed) × seq_page_cost
                = 9346 × 1.0
                = 9346.00
```

```
- total_cost = startup_cost + cpu_run_cost + disk_run_cost
              = 0.00 + 12996.625 + 9346.00
              = 22342.625 ✓
```

-- 5 Simple Seq Scan (with complex subquery filter)

EXPLAIN VERBOSE

```
SELECT i.a
FROM indexed AS i
WHERE i.a <= (SELECT AVG(i.a) FROM indexed AS i);
```


QUERY PLAN

```
Seq Scan on public.indexed i (cost=21846.01..46192.01 rows=333333 width=4)
  Output: i.a
  Filter: ((i.a)::numeric <= $0)
  InitPlan 1 (returns $0)
    -> Aggregate (cost=21846.00..21846.01 rows=1 width=32)
      Output: avg(i_1.a)
      -> Seq Scan on public.indexed i_1 (cost=0.00..19346.00 rows=1000000 width=4)
        Output: i_1.a, i_1.b, i_1.c
```

```
set enable_indexscan = on;
set enable_bitmapscan = on;
```

- Complex predicate  $i.a \leq (\text{SELECT AVG}(i.a) \text{ FROM indexed AS } i)$ :
- $\text{startup\_cost}(i.a \leq (\text{SELECT AVG}(i.a) \text{ FROM indexed AS } i))$ 
  - =  $\text{run\_cost}(\text{SELECT AVG}(i.a) \text{ FROM indexed AS } i)$
  - = 21846.01
- $\text{run\_cost}(i.a \leq (\text{SELECT AVG}(i.a) \text{ FROM indexed AS } i))$ 
  - =  $2 \times \text{cpu\_operator\_cost} \cdot :: \text{numeric}, \cdot \leq \$0$  ( $\$0$  is a constant once InitPlan 1 has been evaluated)
  - =  $2 \times 0.0025$
- $\text{sel}(i.a \leq (\text{SELECT AVG}(i.a) \text{ FROM indexed AS } i)) = 333333 / 1000000 = 0.33$   $\leftarrow$  A arbitrary (1/3)
- $\text{startup\_cost} = \text{startup\_cost}(\text{pred}) + \text{startup\_cost}(\text{expr})$ 
  - =  $21846.01 + 0$
  - = 21846.01 ✓
- $\text{cpu\_run\_cost} = \# \text{rows}(\text{indexed}) \times (\text{cpu\_tuple\_cost} + \text{run\_cost}(\text{pred}))$ 
  - +  $\# \text{rows}(\text{indexed}) \times \text{sel}(\text{pred}) \times \text{run\_cost}(\text{expr})$
  - =  $10^6 \times (0.01 + 2 \times 0.0025)$
  - +  $10^6 \times 0.33 \times 0$
  - = 15000.0
- $\text{disk\_run\_cost} = \# \text{pages}(\text{indexed}) \times \text{seq\_page\_cost}$ 
  - =  $9346 \times 1.0$
  - = 9346.00
- $\text{total\_cost} = \text{startup\_cost} + \text{cpu\_run\_cost} + \text{disk\_run\_cost}$ 
  - =  $21846.01 + 15000.0 + 9346.00$
  - = 46192.01 ✓

## Cost of Index Scan ①



Modeling the cost for an **Index Scan** has to reflect that *two* data structures (heap file & B+Tree) are involved:


<i>idx</i>	<i>in</i>	QUERY PLAN
		Index Scan using indexed_a on indexed i (cost=0.42..443.12 rows=10885 ...)
		Output: (c + '1'::numeric) $\leftarrow$ <i>expr</i>
		Index Cond: (i.a <= 10000) $\leftarrow$ <i>pred</i>
		#rows(out)

The model separately accounts for

1. the B+Tree descent (startup of the **Index Scan**),
2. the index leaf level scan, and
3. heap file access (clustered vs. non-clustered).



## Cost of Index Scan ②





Cost model parameters:

Parameter	Description	Available as...
<code>#rows(<i>in</i>)</code>	# rows (cardinality) of table <i>in</i>	<code>pg_class.reltuples</code>
<code>#pages(<i>in</i>)</code>	# pages in heap file of <i>in</i>	<code>pg_class.relpages</code>
<code>sel(<i>pred</i>)</code>	selectivity of filter <i>pred</i>	see below
<code>h(<i>idx</i>)</code>	height of B+Tree <i>idx</i>	<code>bt_metap(.)</code>
<code>#rows(<i>idx</i>)</code>	# leaf entries in index <i>idx</i>	<code>pg_class.reltuples</code>
<code>#pages(<i>idx</i>)</code>	# pages in leaf level of <i>idx</i>	<code>pg_class.relpages</code>
<code>corr(<i>idx</i>)</code>	$\approx$ clustering factor for index <i>idx</i>	<code>pg_stats.correlation</code>

- `corr(idx)`  $\in \{-1.0, \dots, 1.0\}$  characterizes how much the physical orderings of index leaves and heap file deviate.
  - After `CLUSTER in ON idx`, we have `corr(idx) = 1.0`.

## Cost of Index Scan ③ (B+Tree Descent)



- B+Tree height  $h = \log_{2 \times o}(\#rows(idx))$

⇒ # of key comparisons during B+Tree descent ①:

$$\underbrace{\lceil \log_2(2 \times o) \times h \rceil}_{\text{binary search in inner B+Tree node with fan-out } F = 2 \times o} = \lceil \log_2(\#rows(idx)) \rceil$$

binary search in inner B+Tree  
node with fan-out  $F = 2 \times o$

$$\begin{aligned} \text{startup\_cost} &\stackrel{\text{def}}{=} \text{startup\_cost}(pred) + \text{startup\_cost}(expr) \\ &+ \underbrace{\lceil \log_2(\#rows(idx)) \rceil}_{\text{B+Tree descent}} + \underbrace{(h + 1)}_{\text{: + •}} \times \underbrace{50}_{\text{index node processing}} \times \text{cpu\_operator\_cost} \end{aligned}$$

$$\log_2(2 \times o) \times h$$


$$= \log_2(2 \times o) \times \log_{2 \times o}(\#entries(idx))$$

$$= \frac{\log(2 \times o)}{\log(2)} \times \frac{\log(\#entries(idx))}{\log(2 \times o)}$$


$$= \frac{\log(\#entries(idx))}{\log(2)}$$

$$= \log_2(\#entries(idx))$$

## Cost of Index Scan ④ (Leaf Level Scan)




The index leaf level (sequence set) scan ② incurs CPU as well as I/O cost that contribute to the overall **run\_cost**:



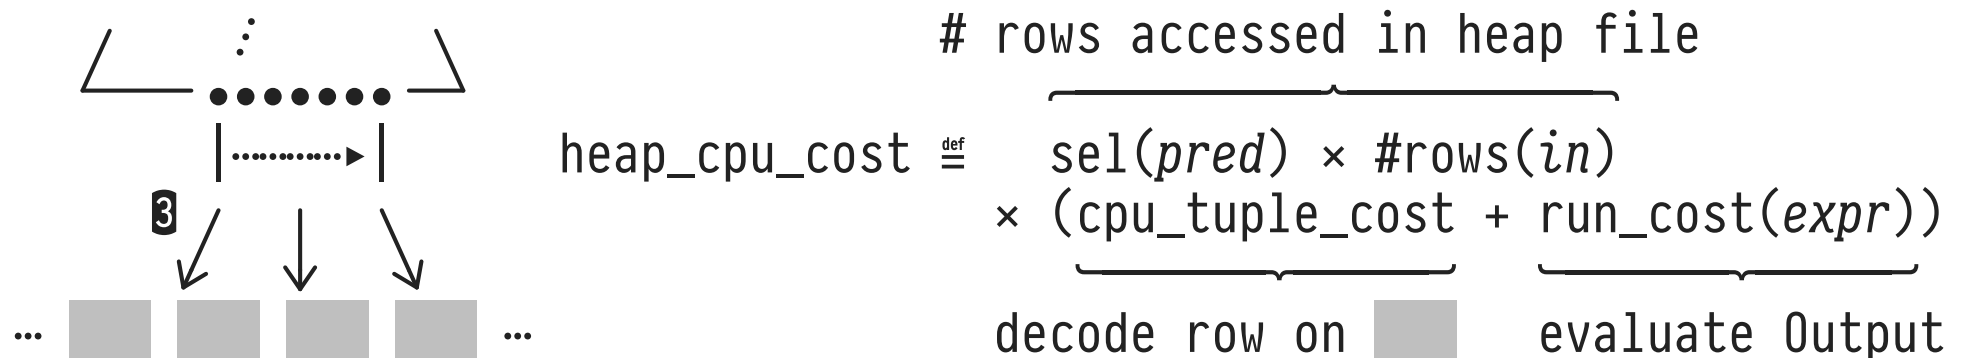
$$\text{index\_cpu\_cost} \stackrel{\text{def}}{=} \underbrace{\text{sel}(\text{pred}) \times \# \text{rows}(\text{idx})}_{\text{\# rows in scanned range}} \times \underbrace{(\text{cpu\_index\_tuple\_cost})}_{\text{decode index leaf entry}} + \underbrace{\text{run\_cost}(\text{pred})}_{\text{evaluate } \leq hi}$$

$$\text{index\_IO\_cost} \stackrel{\text{def}}{=} \underbrace{[\text{sel}(\text{pred}) \times \# \text{pages}(\text{idx})]}_{\text{\# of pages • in scanned range}} \times \underbrace{\text{random\_page\_cost}}_{\text{B+Tree leaves not clustered}}$$

## Cost of Index Scan ⑤ (Heap File Access)



Heap file accesses ③ incur additional CPU and I/O costs (no I/O cost if we perform an **Index Only Scan**):



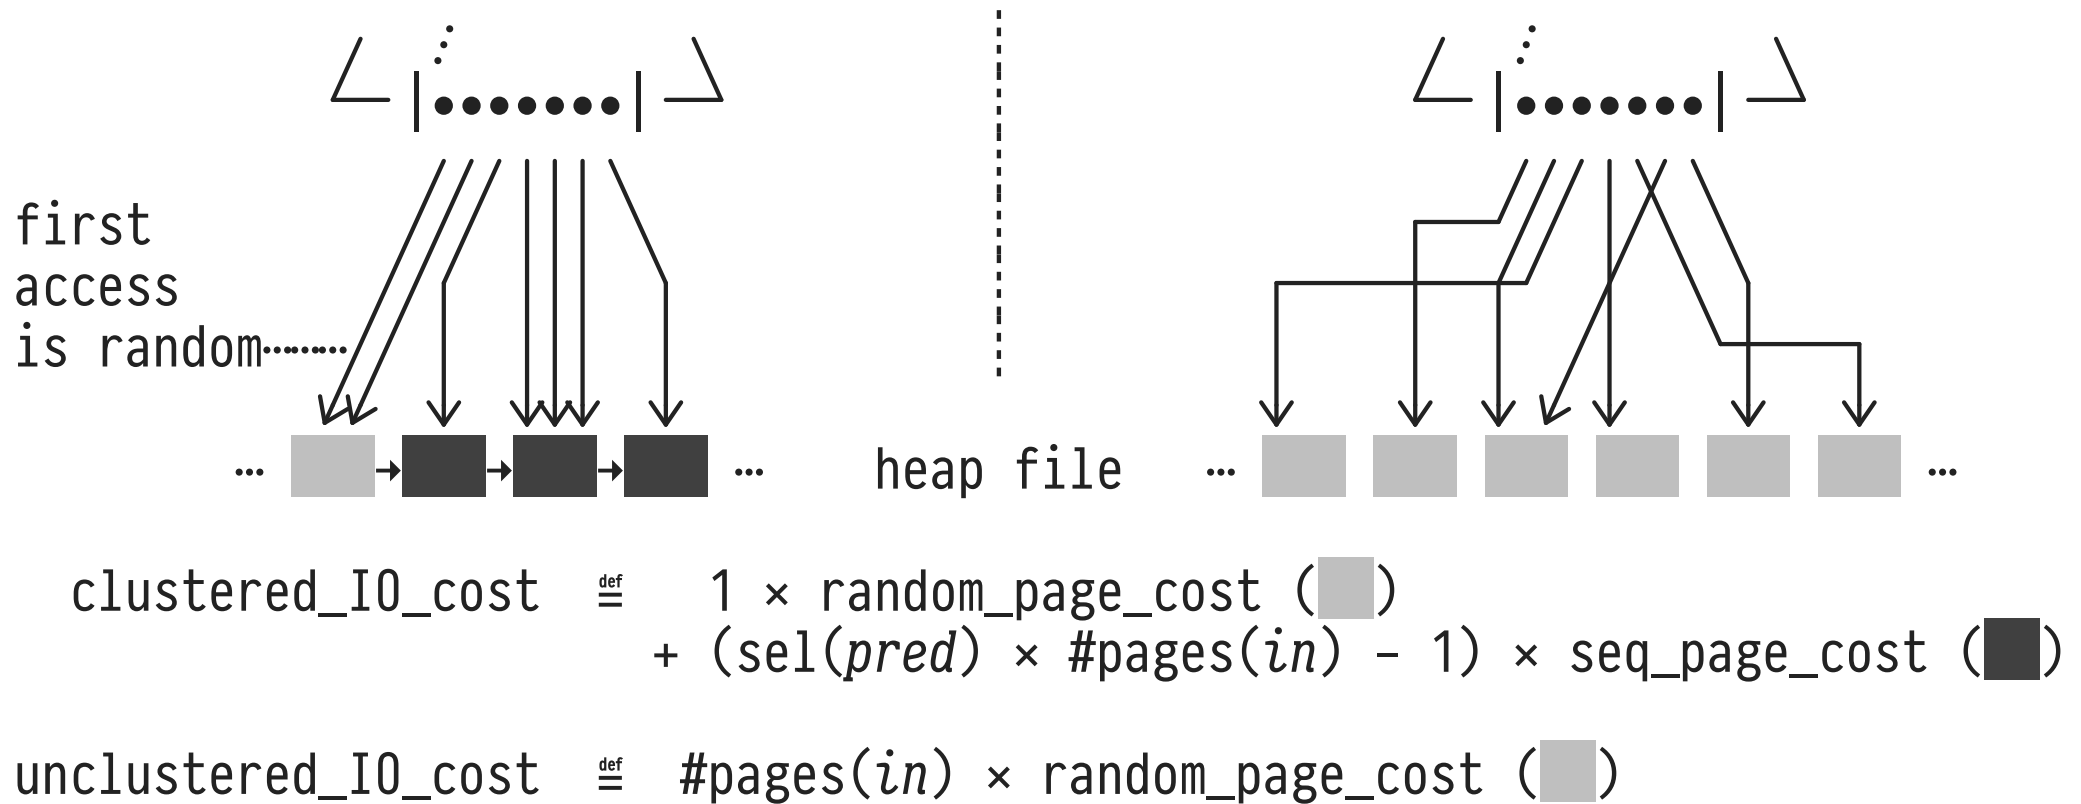
- The more **clustered** the index, the cheaper the heap I/O. Linearly interpolate between the clustered and non-clustered scenarios:

$$\text{heap\_IO\_cost} \stackrel{\text{def}}{=} \text{unclustered\_IO\_cost} + \text{corr}(\text{idx})^2 \times (\text{clustered\_IO\_cost} - \text{unclustered\_IO\_cost})$$

$\uparrow \approx \text{clustering factor} \in \{0, \dots, 1\}$


## Cost of Index Scan ⑥ ([Non-]Clustered Heap File Access)

---



$$\text{total\_cost} \stackrel{\text{def}}{=} \text{startup\_cost} + \text{index\_cpu\_cost} + \text{index\_IO\_cost} + \text{heap\_cpu\_cost} + \text{heap\_IO\_cost}$$

## Index Correlation (Clustering Factor)



Given ordered index  $idx$  over column  $A$  with values  $a_1 \leq a_2 \leq \dots \leq a_n$ , where  $pos(a_i) \in \{1, \dots, n\}$  gives the position of  $a_i$  in the heap file for  $A$ .<sup>6</sup>

- **Index Correlation**  $corr(idx) \in \{-1, \dots, 1\}$  measures how far  $[pos(a_1), \dots, pos(a_n)]$  deviates from  $[1, \dots, n]$ , i.e.,  $idx$ 's clustering degree:

$$corr(idx) = \frac{n \times (\sum_{i=1 \dots n} i \times pos(a_i)) - (\sum_{i=1 \dots n} i)^2}{n \times (\sum_{i=1 \dots n} i \times i) - (\sum_{i=1 \dots n} i)^2}$$

<sup>6</sup> After `CLUSTER <table> USING idx`, we have  $pos(a_i) = i$  and thus  $corr(idx) = 1$ .

Demonstrate the cost derivation for Seq Scan on table `indexed`.

-- ❶ Prepare input table and indexes

```
CREATE INDEX indexed_c ON indexed USING btree (c);
CLUSTER indexed USING indexed_c;
ANALYZE indexed;
```

\d indexed

Table "public.indexed"

Column	Type	Collation	Nullable	Default
a	integer		not null	
b	text			
c	numeric(3,2)			

Indexes:

```
"indexed_a" PRIMARY KEY, btree (a)
"indexed_c" btree (c) CLUSTER
```

-- ❷ Obtain meta data about table indexed

```
SELECT relname, reltuples AS "#rows(.)", relpages AS "#pages(.)"
FROM   pg_class
WHERE  relname LIKE 'indexed%';
```

relname	#rows(.)	#pages(.)
indexed	1e+06	9343
indexed_a	1e+06	2745
indexed_c	1e+06	2745

```
SELECT correlation AS "corr(indexed_a)"
FROM   pg_stats
WHERE  tablename = 'indexed' AND attname = 'a';
```

corr(indexed_a)
0.00518881

```
SELECT level AS "h(indexed_a)"
FROM   bt_metap('indexed_a');
```



h(indexed_a)
2

```
-- 8 Perform a index range scan over a non-clustered index
-- (cf. with Seq Scan query 5 above which had cost=0.00..22342.17)
```

```
set enable_bitmapscan = off;
set enable_seqscan = off;
```

```
EXPLAIN VERBOSE
SELECT i.c * 2 + 1
FROM indexed AS i
WHERE i.a <= 100000;
```

QUERY PLAN
Index Scan using indexed_a on public.indexed i (cost=0.42..40779.96 rows=101712 width=32) Output: ((c * '2'::numeric) + '1'::numeric) Index Cond: (i.a <= 100000)

```
- sel(pred): sel(i.a <= 100000) = 101712 / 1000000 = 0.101712

- startup_cost = startup_cost(pred) + startup_cost(expr)
                + ([log2(#rows(idx))] + (h + 1) × 50) × cpu_operator_cost
                = 0 + 0
                + ([log2(106)] + (2 + 1) × 50) × 0.0025
                = 0.42 ✓

- index_cpu_cost = sel(pred) × #rows(indexed_a)
                  × (cpu_index_tuple_cost + run_cost(pred))
                  = 0.101712 × 106
                  × (0.005 + 1 × 0.0025)   ◀ run_cost(pred) = 1 × cpu_operator_cost: • <= 100000
                  = 762.84

- index_IO_cost = [sel(pred) × #pages(indexed_a)] × random_page_cost
                = [0.101712 × 2745] × 4.0
                = 1120

- heap_cpu_cost = sel(pred) × #rows(indexed)
                 × (cpu_tuple_cost + run_cost(expr))
                 = 0.101712 × 106
                 × (0.01 + 2 × 0.0025)   ◀ run_cost(expr) = 2 × cpu_operator_cost: • * 2, • + 1
```

= 1525.68

```
- clustered_I0_cost = 1 × random_page_cost + (sel(pred) * #pages(indexed) - 1) × seq_page_cost
                  = 1 × 4.0 + (0.101712 × 9343 - 1) × 1.0
                  = 953.295

- unclustered_I0_cost = #pages(indexed) × random_page_cost
                    = 9343 × 4.0
                    = 37372.0

- corr(indexed_a) = 0.00518881

- heap_I0_cost = unclustered_I0_cost
               + corr(indexed_a)2 × (clustered_I0_cost - unclustered_I0_cost)
               = 37372.0
               + 0.005188812 × (953.295 - 37372.0)
               = 37371.0194

- total_cost = startup_cost + index_cpu+cost + index_I0_cost
              + heap_cpu_cost + heap_I0_cost
              = 0.42 + 762.84 + 1120
                + 1525.68 + 37371.0194
              = 40779.9594 ✓
```

```
-- ❹ Perform a index range scan over a clustered index
-- (cf. with Seq Scan query ❸ above which had cost=0.00..22342.17)
```

```
CLUSTER indexed USING indexed_a;
ANALYZE indexed;
```

```
SELECT correlation AS "corr(indexed_a)"
FROM   pg_stats
WHERE  tablename = 'indexed' AND attname = 'a';
```

corr(indexed_a)
1

```
EXPLAIN VERBOSE
SELECT i.c * 2 + 1
FROM indexed AS i
WHERE i.a <= 100000;
```

# QUERY PLAN

Index Scan using indexed\_a on public.indexed i (cost=0.42..4299.33 rows=100218 width=32)  
Output: ((c \* '2'::numeric) + '1'::numeric)  
Index Cond: (i.a <= 100000)

```

- sel(pred): sel(i.a <= 100000) = 100218 / 1000000 = 0.100218

- startup_cost = startup_cost(pred) + startup_cost(expr)
                + ([log2(#rows(idx))] + (h + 1) × 50) × cpu_operator_cost
                = 0 + 0
                + ([log2(106)] + (2 + 1) × 50) × 0.0025
                = 0.42 ✓

- index_cpu_cost = sel(pred) × #rows(indexed_a)
                  × (cpu_index_tuple_cost + run_cost(pred))
                  = 0.100218 × 106
                  × (0.005 + 1 × 0.0025)   ← run_cost(pred) = 1 × cpu_operator_cost: • <= 100000
                  = 751.63

- index_IO_cost = [sel(pred) × #pages(indexed_a)] × random_page_cost
                 = [0.100218 × 2745] × 4.0
                 = 1104

- heap_cpu_cost = sel(pred) × #rows(indexed)
                 × (cpu_tuple_cost + run_cost(expr))
                 = 0.100218 × 106
                 × (0.01 + 2 × 0.0025)   ← run_cost(expr) = 2 × cpu_operator_cost: • * 2, • + 1
                 = 1503.27

- clustered_IO_cost = 1 × random_page_cost + (sel(pred) × #pages(indexed) - 1) × seq_page_cost
                   = 1 × 4.0 + (0.100218 × 9343 - 1) × 1.0
                   = 939.336

- unclustered_IO_cost = #pages(indexed) × random_page_cost
                     = 9343 × 4.0
                     = 37372.0

- corr(indexed_a) = 1.0

- heap_IO_cost = unclustered_IO_cost
                + corr(indexed_a)2 × (clustered_IO_cost - unclustered_IO_cost)
                = 37372.0
                + 1.02 × (939.336 - 37372.0)
                = 939.336

```

```

- total_cost = startup_cost + index_cpu_cost + index_IO_cost
               + heap_cpu_cost + heap_IO_cost
= 0.42 + 751.63 + 1104
      + 1503.27 + 939.336
= 4298.656 ✓

```

```

-- 5 Perform index-ONLY scan over a clustered index

```

```

-- make sure, dead rows are removed (visibility map update)
VACUUM;

```

```

EXPLAIN VERBOSE
SELECT i.a * 2 + 1
FROM indexed AS i
WHERE i.a <= 100000;

```

QUERY PLAN
Index Only Scan using indexed_a on public.indexed i (cost=0.42..3359.33 rows=100218 width=4) Output: ((a * 2) + 1) Index Cond: (i.a <= 100000)

```

- sel(pred): sel(i.a <= 100000) = 100218 / 1000000 = 0.100218

- startup_cost = startup_cost(pred) + startup_cost(expr)
                 + ([log2(#rows(idx))] + (h + 1) × 50) × cpu_operator_cost
= 0 + 0
+ ([log2(106)] + (2 + 1) × 50) × 0.0025
= 0.42 ✓

- index_cpu_cost = sel(pred) × #rows(indexed_a)
                  × (cpu_index_tuple_cost + run_cost(pred))
= 0.100218 × 106
× (0.005 + 1 × 0.0025)   ← run_cost(pred) = 1 × cpu_operator_cost: • <= 100000
= 751.63

- index_IO_cost = [sel(pred) × #pages(indexed_a)] × random_page_cost
= [0.100218 × 2745] × 4.0
= 1104

```

```

-- We do not access the heap file to process the rows but still need to
-- process the rows read off the index leaves, thus heap_cpu_cost > 0 even

```

```

-- for an Index Only Scan:

- heap_cpu_cost = sel(pred) × #rows(indexed)
                  × (cpu_tuple_cost + run_cost(expr))
                  = 0.100218 × 106
                  × (0.01 + 2 × 0.0025)  ← run_cost(expr) = 2 × cpu_operator_cost: . * 2, . + 1
                  = 1503.27

- heap_IO_cost  = 0 (Index Only Scan)

- total_cost = startup_cost + index_cpu+cost + index_IO_cost
               + heap_cpu_cost + heap_IO_cost
               = 0.42 + 751.63 + 1104
                 + 1503.27 + 0
               = 3359.32 ✓

```

**TODO:**

- Explain selectivity estimation in PostgreSQL [mcv, histograms, correlation and CREATE STATISTICS?]
- Explain cost model for a join operator (see <http://www.interdb.jp/pg/pgsql103.html> for cost models for (Index) Nested Loop Join, Merge Join)